

Cross-Disease Graph Diagnosis: A Clinical Benchmark of CNN and GNN Architectures and a Falsifiable Test of Cross-Disease Patient Bridges

MEGA-27 Item 10 - Deep-Learning Diagnosis Suite. September 24, 2026.

Udita Phookan

Working revision, October 7, 2026. Locked predictive gates failed. Dataset independence, scientific-tool breadth, novelty and matched benchmark superiority remain open. The PDF uses Liberation Serif because actual Times New Roman is unavailable. Page length includes tables and appendices, not an equivalent amount of new research.

Abstract

We benchmark MLP, 1D-CNN, GCN and GAT diagnosis models against strong classical baselines on real clinical datasets (recorded sources, independent census unverified), with bootstrap confidence intervals and calibration metrics, and we test a falsifiable discovery claim: cross-disease patient bridges - edges of a unified patient-similarity graph spanning diseases, built on disease-agnostic distributional fingerprints - carry transferable diagnostic signal. All claims are reported with seed variance; negative results are preserved.

Introduction

Computer-aided diagnosis is one of the oldest promises of machine learning, yet the published benchmark culture around small clinical datasets is methodologically fragile. The Wisconsin Diagnostic Breast Cancer (WDBC), Cleveland heart disease, Pima diabetes and Parkinson's vocal-measurement cohorts are widely reused benchmarks. This paper does not establish a systematic prevalence estimate for leakage or headline-score inflation in the literature. Non-nested selection and test-set peeking are risks to check in a matched comparator, rather than accusations inferred from a high reported score.

This project asks three questions under a deliberately strict protocol (held-out 25 percent test set, five seeds, bootstrap confidence intervals, calibration metrics, hermetic tests for every computation):

1. Floor: what do modern deep architectures (MLP, 1D-CNN, GCN, GAT) actually achieve on real clinical cohorts when evaluated honestly, and how do they compare to strong classical baselines (regularized logistic regression, random forests)?
2. Discovery: does a unified patient-similarity graph spanning multiple diseases contain transferable diagnostic signal? We formalize this as the bridge hypothesis: cross-disease edges of a shared graph, built on

disease-agnostic distributional fingerprints, improve per-disease diagnosis in a paired ablation across seeds.

3. Label efficiency: can a transductive graph model trained on a fraction of labels match or beat a full-label tabular model? This is the regime that matters clinically, where labeled data is the bottleneck.

All code, every raw number behind every table, and all negative results are preserved in the project repository.

Methods

Datasets. Twenty-six named core/panel result entries are recorded (four core and twenty-two panel files). Their source routes are UCI or cited public mirrors. These files are not an independently deduplicated dataset census. HTTP success alone does not verify source identity, labels, independence or current availability; the historical URL-check claim is not certified. The four core cohorts (WDBC, Cleveland heart disease, Pima diabetes, Parkinson's) receive the full six-model, five-seed treatment; a further twenty-two datasets (four heart-disease sites, two thyroid cohorts, two breast-cancer cohorts, hepatitis, SPECTF cardiac scintigraphy, Haberman survival, lymphography, dermatology, mammographic masses, echocardiogram, arrhythmia, cervical-cancer risk factors, Messidor diabetic retinopathy, fertility, HCV, Indian liver patient dataset, early-stage diabetes risk) receive a four-model, three-seed panel treatment. Missing values are median-imputed; categorical variables are factorized; labels are binarized to the clinically standard positive class (documented per dataset in the repository loaders).

Models. The MLP is a two-hidden-layer ReLU network trained with Adam and class-weighted binary cross-entropy. The 1D-CNN treats the standardized feature vector as a one-channel sequence with two convolutional blocks. The GCN and GAT operate on a patient-similarity graph: nodes are patients, edge weights come from a Gaussian kernel over k-nearest neighbors with a self-tuned bandwidth (Eq. 1), and propagation uses the symmetrically normalized adjacency (Eq. 2). Graph models are trained transductively: the graph includes test patients' features but never their labels.

Unified cross-disease graph. Each disease's encoder maps its native feature space into a shared 16-dimensional latent space. Every patient is additionally described by a 10-dimensional distributional fingerprint (Eq. 4) - moments and quantiles of the patient's own feature vector - which discards feature identity but is not thereby comparable across diseases. Raw moments and quantiles still depend on units, feature sets and cohort composition. Only the standardized moments have the restricted common rescaling invariance stated in Derivation 8. The unified graph is the union of the latent

kNN graph and the fingerprint kNN graph; a two-layer GCN trunk with one output head per disease is trained on the summed per-disease losses (Eq. 6). The ablated variant removes exactly the cross-disease edges.

Evaluation. Every run reports ROC AUC (Mann-Whitney form, Eq. 10), balanced accuracy, Brier score, and expected calibration error (Eq. 7), with percentile bootstrap 95 percent confidence intervals (Eq. 8). The bridge effect per disease is the paired per-seed AUC difference between the full and ablated graphs (Eq. 9); a hypothetical conditional sign-test formula is supplied (Proposition 1), but its independence and fair-sign assumptions are not established for seeds sharing patients. Bridge differences are descriptive. Label-efficiency experiments subsample training labels at 10, 25, 50 and 100 percent, five seeds, comparing transductive GCN against MLP trained on the same labeled subset.

Reproducibility. The full suite ships with hermetic pytest suites (28 tests) that run without network access; live dataset downloads happen only through a single documented fetch function with on-disk caching. Seeds, splits and every hyperparameter are recorded in the committed result JSONs.

Mathematical formulation

(1) Gaussian kNN kernel, self-tuned bandwidth: $A_{ij} = \exp(-\|x_i - x_j\|^2 / (2 \sigma^2))$, $\sigma^2 = \text{median kNN squared distance}$ (Zelnik-Manor & Perona).

(2) Symmetric GCN normalization: $A_{\text{hat}} = D^{-1/2} (A + I) D^{-1/2}$ (Kipf & Welling 2017).

(3) GCN layer: $H^{l+1} = \tanh(A_{\text{hat}} H^l W^l)$. Derivation: first-order Chebyshev approximation of spectral graph convolution with $\lambda_{\text{max}} \approx 2$ and renormalization trick.

(4) Distributional fingerprint $\phi(x) = (\mu, \sigma, \min, \max, q_{25}, q_{50}, q_{75}, m_3, m_4, |z|_{1/d}, m_r = \text{mean } z_j^r, z_j = (x_j - \mu)/\sigma)$ - feature identities discarded; cross-disease comparability unproven.

(5) Class-weighted BCE: $L = -(1/n) \sum_i w_{y_i} [y_i \log s(f_i) + (1 - y_i) \log(1 - s(f_i))]$, $w_1 = n_-/n_+$.

(6) Multi-task transductive objective: $L_{\text{MT}} = \sum_d L^{(d)}(T_d)$; $\text{grad } L_{\text{MT}} = \sum_d \text{grad } L^{(d)}$.

(7) ECE = $\sum_b (|S_b|/n) |\text{acc}(S_b) - \text{conf}(S_b)|$ over 10 equal-width bins.

(8) Percentile bootstrap $\text{CI}_{95} = [Q_{0.025}(\theta^*_b), Q_{0.975}(\theta^*_b)]$, $B = 1000$ resamples.

(9) Bridge ablation: $\Delta^{(d)} = (1/S) \sum_s (\text{AUC}_{\text{full}} - \text{AUC}_{\text{no-bridge}})$; conditional one-sided sign-test $p = 2^{-S} \sum_{k \geq k_0} C(S, k)$; independence and sign symmetry are unverified.

(10) Mann-Whitney AUC: $AUC = P(f(x+) > f(x-)) + 0.5 P(f(x+) = f(x-))$.

(11) Bridge homophily $h = |\{(i,j) \text{ in } E_{\text{cross}} : y_i = y_j\}| / |E_{\text{cross}}|$.

Datasets

Dataset	n	features	Citation
cleveland	303	13	Detrano et al. (1989), UCI ML Repository
ann_thyroid	7200	21	Quinlan, UCI (ANN thyroid)
arrhythmia	452	279	Guvenir et al. (1997), UCI
cervical	858	35	Fernandes et al. (2017), UCI
dermatology	366	34	Ilter & Guvenir (1998), UCI
diabetes_early_risk	520	16	Islam et al. (2020), UCI
echocardiogram	131	9	UCI ML Repository
fertility	100	9	Gil et al. (2012), UCI
haberman	306	3	Haberman (1976), UCI
hcv	615	12	Lichtinghagen et al. / Hoffmann et al. (2018), UCI
heart_hungarian	294	13	Detrano et al. (1989), UCI (hungarian site)
heart_switzerland	123	13	Detrano et al. (1989), UCI (switzerland site)
heart_va	200	13	Detrano et al. (1989), UCI (va site)
hepatitis	155	19	UCI ML Repository
ilpd	583	10	Ramana & Venkateswarlu, UCI
lymphography	148	18	UCI ML Repository
mammographic	961	5	Elter et al. (2007), UCI
messidor	1151	19	Antal & Hajdu (2014), UCI
new_thyroid	215	5	UCI ML Repository
spectf	267	44	Kurgan et al. (2001), UCI
statlog_heart	270	13	StatLog / UCI
wisconsin_original	699	9	Wolberg & Mangasarian (1990), UCI
wpbc	198	34	Wolberg et al., UCI

			(prognostic breast cancer)
parkinsons	195	22	Little et al. (2007), UCI ML Repository
pima	768	8	Smith et al. (1988), NIDDK / UCI ML Repository
wdbc	569	30	Street, Wolberg & Mangasarian (1993), UCI ML Repository

Results

Each subsection reports the committed result file for one dataset: per-model ROC AUC (mean $\pm$ std across seeds), balanced accuracy, Brier score and ECE. The best model is named; all numbers trace to results/*.json in the repository.

Dataset: cleveland

Model	ROC AUC	Balanced acc	Brier	ECE
mlp	0.8820 $\pm$ 0.0268	0.8020	0.1458	0.1234
cnn1d	0.8838 $\pm$ 0.0149	0.7861	0.1425	0.1115
gcn	0.8808 $\pm$ 0.0209	0.7966	0.1400	0.0926
gat	0.8728 $\pm$ 0.0255	0.7925	0.1449	0.1120
logreg	0.8902 $\pm$ 0.0304	0.8109	0.1326	0.0984
random_forest	0.8846 $\pm$ 0.0288	0.8019	0.1356	0.0913

Dataset: ann_thyroid

Model	ROC AUC	Balanced acc	Brier	ECE
mlp	0.9857 $\pm$ 0.0093	0.9624	0.0120	0.0117
gcn	0.9314 $\pm$ 0.0198	0.8620	0.0636	0.1245
logreg	0.9854 $\pm$ 0.0012	0.9729	0.0292	0.0769
random_forest	0.9994 $\pm$ 0.0002	0.9914	0.0053	0.0156

Dataset: arrhythmia

Model	ROC AUC	Balanced acc	Brier	ECE
mlp	0.8195 $\pm$ 0.0372	0.7519	0.1711	0.0730
gcn	0.7701 $\pm$ 0.0269	0.7075	0.1960	0.0976

logreg	0.7713 +- 0.0150	0.7086	0.2247	0.1977
random_forest	0.8724 +- 0.0205	0.8182	0.1554	0.1313

Dataset: cervical

Model	ROC AUC	Balanced acc	Brier	ECE
mlp	0.9362 +- 0.0252	0.8976	0.0418	0.0453
gcn	0.9132 +- 0.0414	0.8552	0.0549	0.0716
logreg	0.9294 +- 0.0363	0.8815	0.0519	0.0700
random_forest	0.9480 +- 0.0278	0.6941	0.0348	0.0314

Dataset: dermatology

Model	ROC AUC	Balanced acc	Brier	ECE
mlp	1.0000 +- 0.0000	1.0000	0.0000	0.0012
gcn	1.0000 +- 0.0000	1.0000	0.0002	0.0028
logreg	1.0000 +- 0.0000	1.0000	0.0013	0.0117
random_forest	1.0000 +- 0.0000	1.0000	0.0059	0.0346

Dataset: diabetes_early_risk

Model	ROC AUC	Balanced acc	Brier	ECE
mlp	0.9927 +- 0.0024	0.9521	0.0412	0.0541
gcn	0.9882 +- 0.0032	0.9300	0.0484	0.0545
logreg	0.9668 +- 0.0102	0.8975	0.0728	0.0621
random_forest	0.9975 +- 0.0012	0.9679	0.0257	0.0600

Dataset: echocardiogram

Model	ROC AUC	Balanced acc	Brier	ECE
mlp	0.8251 +- 0.0439	0.7197	0.1602	0.1459
gcn	0.8127 +- 0.0594	0.7197	0.1797	0.1797
logreg	0.8333 +- 0.0128	0.6742	0.1634	0.2079
random_forest	0.8251 +- 0.0364	0.7348	0.1378	0.1418

Dataset: fertility

Model	ROC AUC	Balanced acc	Brier	ECE
mlp	0.6591 +- 0.0528	0.5354	0.2202	0.3399
gcn	0.4621 +- 0.1269	0.5455	0.2032	0.3081
logreg	0.4672 +- 0.1090	0.5480	0.2797	0.3512
random_forest	0.4798 +- 0.1149	0.4773	0.1350	0.1257

Dataset: haberman

Model	ROC AUC	Balanced acc	Brier	ECE
mlp	0.6625 +- 0.0401	0.6351	0.2192	0.1915
gcn	0.6336 +- 0.0149	0.6190	0.2233	0.1939
logreg	0.6452 +- 0.0702	0.6417	0.2180	0.2185
random_forest	0.6783 +- 0.0055	0.5613	0.2031	0.1371

Dataset: hcv

Model	ROC AUC	Balanced acc	Brier	ECE
mlp	0.9672 +- 0.0107	0.8804	0.0374	0.0356
gcn	0.9535 +- 0.0284	0.8774	0.0476	0.0534
logreg	0.9405 +- 0.0327	0.8904	0.0541	0.0834
random_forest	0.9670 +- 0.0208	0.8296	0.0333	0.0532

Dataset: heart_hungarian

Model	ROC AUC	Balanced acc	Brier	ECE
mlp	0.9176 +- 0.0062	0.8129	0.1119	0.1002
gcn	0.9059 +- 0.0211	0.8122	0.1127	0.1058
logreg	0.9184 +- 0.0157	0.8286	0.1059	0.1013
random_forest	0.9026 +- 0.0318	0.7793	0.1243	0.0908

Dataset: heart_switzerland

Model	ROC AUC	Balanced acc	Brier	ECE
mlp	0.6954 +- 0.2263	0.6925	0.1050	0.1379
gcn	0.7759 +- 0.0000	0.6466	0.1322	0.2040

	0.1623			
logreg	0.8678 +- 0.0849	0.6580	0.1211	0.2018
random_forest	0.6868 +- 0.0935	0.4943	0.0615	0.0617

Dataset: heart_va

Model	ROC AUC	Balanced acc	Brier	ECE
mlp	0.7027 +- 0.0465	0.5970	0.2326	0.2511
gcn	0.5468 +- 0.1163	0.5922	0.2557	0.2735
logreg	0.6999 +- 0.0466	0.6923	0.2258	0.2412
random_forest	0.7190 +- 0.0311	0.5658	0.1729	0.1245

Dataset: hepatitis

Model	ROC AUC	Balanced acc	Brier	ECE
mlp	0.8777 +- 0.0575	0.8353	0.1392	0.2218
gcn	0.9315 +- 0.0456	0.8622	0.1289	0.2223
logreg	0.8978 +- 0.0363	0.8259	0.1050	0.1363
random_forest	0.9288 +- 0.0580	0.7964	0.0817	0.1345

Dataset: ilpd

Model	ROC AUC	Balanced acc	Brier	ECE
mlp	0.7538 +- 0.0479	0.7019	0.2084	0.1842
gcn	0.7116 +- 0.0385	0.6049	0.2140	0.1578
logreg	0.7511 +- 0.0482	0.6916	0.2014	0.1696
random_forest	0.7647 +- 0.0403	0.6017	0.1698	0.0671

Dataset: lymphography

Model	ROC AUC	Balanced acc	Brier	ECE
mlp	0.8978 +- 0.0388	0.8457	0.1192	0.1647
gcn	0.8492 +- 0.0258	0.7852	0.1444	0.1595
logreg	0.8403 +- 0.0501	0.7956	0.1515	0.1593
random_forest	0.9107 +- 0.0462	0.8175	0.1228	0.1362

Dataset: mammographic

Model	ROC AUC	Balanced acc	Brier	ECE
mlp	0.9019 +- 0.0057	0.8172	0.1284	0.0675
gcn	0.8956 +- 0.0110	0.8219	0.1275	0.0597
logreg	0.8976 +- 0.0125	0.8195	0.1265	0.0411
random_forest	0.8687 +- 0.0144	0.7778	0.1560	0.1097

Dataset: messidor

Model	ROC AUC	Balanced acc	Brier	ECE
mlp	0.7930 +- 0.0209	0.7194	0.1861	0.0601
gcn	0.7075 +- 0.0142	0.6483	0.2166	0.0675
logreg	0.7930 +- 0.0108	0.7133	0.1833	0.0577
random_forest	0.7311 +- 0.0044	0.6628	0.2080	0.0532

Dataset: new_thyroid

Model	ROC AUC	Balanced acc	Brier	ECE
mlp	0.9984 +- 0.0013	0.9688	0.0165	0.0528
gcn	0.9973 +- 0.0021	0.9644	0.0217	0.0531
logreg	0.9364 +- 0.0295	0.8618	0.0983	0.1435
random_forest	0.9995 +- 0.0008	0.9688	0.0213	0.0700

Dataset: spectf

Model	ROC AUC	Balanced acc	Brier	ECE
mlp	0.7960 +- 0.1029	0.7233	0.1748	0.1852
gcn	0.8082 +- 0.0710	0.6925	0.1581	0.1510
logreg	0.8392 +- 0.0557	0.7558	0.1724	0.1765
random_forest	0.8495 +- 0.0466	0.5907	0.1269	0.0810

Dataset: statlog_heart

Model	ROC AUC	Balanced acc	Brier	ECE
mlp	0.9058 +- 0.0069	0.8208	0.1248	0.0910
gcn	0.8991 +- 0.0069	0.8254	0.1226	0.1175

	0.0126			
logreg	0.9038 +- 0.0088	0.8319	0.1207	0.1038
random_forest	0.8917 +- 0.0124	0.8406	0.1304	0.1128

Dataset: wisconsin_original

Model	ROC AUC	Balanced acc	Brier	ECE
mlp	0.9943 +- 0.0003	0.9741	0.0252	0.0403
gcn	0.9955 +- 0.0014	0.9769	0.0229	0.0313
logreg	0.9953 +- 0.0007	0.9687	0.0243	0.0317
random_forest	0.9923 +- 0.0034	0.9686	0.0277	0.0382

Dataset: wpbc

Model	ROC AUC	Balanced acc	Brier	ECE
mlp	0.7039 +- 0.0721	0.6499	0.2097	0.2252
gcn	0.5885 +- 0.0300	0.5395	0.2131	0.1919
logreg	0.7346 +- 0.0224	0.6345	0.2174	0.2240
random_forest	0.6857 +- 0.0232	0.5658	0.1619	0.1036

Dataset: parkinsons

Model	ROC AUC	Balanced acc	Brier	ECE
mlp	0.9635 +- 0.0241	0.8984	0.0783	0.0982
cnn1d	0.9324 +- 0.0442	0.8205	0.0977	0.1141
gcn	0.9595 +- 0.0234	0.8793	0.0802	0.1038
gat	0.8613 +- 0.1061	0.7777	0.1472	0.1581
logreg	0.8982 +- 0.0386	0.7671	0.1356	0.1685
random_forest	0.9484 +- 0.0179	0.8144	0.0801	0.1103

Dataset: pima

Model	ROC AUC	Balanced acc	Brier	ECE
mlp	0.8239 +- 0.0059	0.7418	0.1696	0.0935
cnn1d	0.8205 +- 0.0220	0.7281	0.1702	0.0926

gcn	0.8208 +- 0.0218	0.7497	0.1745	0.1024
gat	0.8175 +- 0.0182	0.7363	0.1749	0.1003
logreg	0.8294 +- 0.0099	0.7551	0.1715	0.1166
random_forest	0.8134 +- 0.0156	0.7277	0.1652	0.0634

Dataset: wdbc

Model	ROC AUC	Balanced acc	Brier	ECE
mlp	0.9917 +- 0.0072	0.9736	0.0207	0.0274
cnn1d	0.9914 +- 0.0050	0.9556	0.0341	0.0413
gcn	0.9902 +- 0.0073	0.9529	0.0300	0.0307
gat	0.9885 +- 0.0085	0.9574	0.0292	0.0275
logreg	0.9929 +- 0.0050	0.9691	0.0224	0.0249
random_forest	0.9899 +- 0.0069	0.9570	0.0315	0.0500

Discovery experiment: cross-disease bridges

Disease	AUC full	AUC ablated	Delta	Delta per seed
wdbc	0.9894	0.9888	+0.0006	[0.004, 0.0006, 0.0019, -0.0013, -0.0021]
cleveland	0.8308	0.8270	+0.0038	[-0.0105, 0.0021, 0.0028, 0.0251, -0.0007]
pima	0.7694	0.7607	+0.0086	[0.0294, 0.0006, 0.0134, 0.0048, -0.005]
parkinsons	0.9009	0.9027	-0.0018	[0.0315, 0.0203, -0.0428, 0.0135, -0.0315]

Stable bridges (seed recurrence $\geq 4/5$): [{"pair": "wdbc<->cleveland", "seed_recurrence": 5, "endpoint_label_counts": {"1-0": 387, "1-1": 1237, "0-0": 829, "0-1": 380}}, {"pair": "wdbc<->pima", "seed_recurrence": 5, "endpoint_label_counts": {"1-0": 1104, "1-1": 2333, "0-0": 3575, "0-1": 1092}}, {"pair": "wdbc<->parkinsons", "seed_recurrence": 5, "endpoint_label_counts": {"1-1": 757, "0-1": 1249, "0-0": 68, "1-0": 24}}, {"pair": "cleveland<->pima", "seed_recurrence": 5, "endpoint_label_counts": {"0-0": 1793, "0-1": 657, "1-1": 1029, "1-0": 1565}}, {"pair": "cleveland<->parkinsons", "seed_recurrence": 5, "endpoint_label_counts": {"1-1": 579, "0-1": 823, "0-0": 167, "1-0": 171}}, {"pair": "pima<->parkinsons", "seed_recurrence": 5, "endpoint_label_counts": {"1-1": 764, "1-0": 212, "0-1": 1094, "0-0": 255}}]

Discovery experiment: cross-disease bridges

Disease	AUC full	AUC ablated	Delta	Delta per seed
wdbc	0.9897	0.9893	+0.0004	[0.0023, -0.0002, 0.0004, -0.0053, 0.0049]
cleveland	0.8273	0.8171	+0.0102	[0.0118, 0.0195, -0.0007, 0.0118, 0.0084]
pima	0.7621	0.7555	+0.0066	[0.0045, 0.0032, 0.0208, 0.0005, 0.0041]
parkinsons	0.9248	0.9108	+0.0140	[0.0293, 0.0473, 0.0, -0.0045, -0.0023]

Stable bridges (seed recurrence $\geq 4/5$): [{"pair": "wdbc<->cleveland", "seed_recurrence": 5, "endpoint_label_counts": {"1-0": 379, "1-1": 1230, "0-0": 821, "0-1": 371}}, {"pair": "wdbc<->pima", "seed_recurrence": 5, "endpoint_label_counts": {"1-0": 1091, "1-1": 2321, "0-0": 3549, "0-1": 1081}}, {"pair": "wdbc<->parkinsons", "seed_recurrence": 5, "endpoint_label_counts": {"1-1": 733, "0-1": 1221, "0-0": 57, "1-0": 15}}, {"pair": "cleveland<->pima", "seed_recurrence": 5, "endpoint_label_counts": {"0-0": 1818, "0-1": 670, "1-1": 1039, "1-0": 1585}}, {"pair": "cleveland<->parkinsons", "seed_recurrence": 5, "endpoint_label_counts": {"1-1": 546, "0-1": 804, "0-0": 155, "1-0": 154}}, {"pair": "pima<->parkinsons", "seed_recurrence": 5, "endpoint_label_counts": {"1-1": 759, "1-0": 205, "0-1": 1101, "0-0": 255}}]

External tools

Tool	Version	Used for
PyTorch	2.14.0+cpu	MLP/CNN1D/GCN/GAT model implementations and training loops
scikit-learn	1.7.2	logistic regression + random forest baselines, metric implementations (roc_auc, brier, f1)
NumPy	2.2.6	all numerical computation, kNN graph construction
SciPy	1.15.3	statistical utilities
NetworkX	latest	graph inspection/validation of patient-similarity graphs
Matplotlib	3.10.9	figures for the paper
Biopython	latest	NCBI/Entrez citation and dataset verification
pytest	latest	hermetic test suites (28 tests)
UCI ML Repository	live	25 clinical datasets, direct-download verified
GitHub (jbrownlee/Datasets)	live	Pima diabetes dataset mirror
python-docx	1.2.0	paper DOCX generation (Times New Roman)

git/GitHub	2.34.1	reproducible artifact versioning
NCBI GEO FTP archive	live	GEO series-matrix downloads for 25-accession transcriptomic arm
OpenML REST API	v1	dataset catalogue query (6000-entry list) and accession-level curation for the OpenML arm
scikit-learn fetch_openml	1.7.2	programmatic fetch of 30 OpenML datasets with caching
OpenML file download endpoint	v1	direct ARFF retrieval + manual parse for heart-failure (did 43008) when default-target metadata was missing
Zenodo	record 10519652	MedMNIST v2 npz direct downloads with resume/integrity validation
MedMNIST package	latest	subset task metadata (INFO) for the 12-subset image arm
pandas	latest	categorical/mixed-type encoding in the OpenML ARFF loader
curl	system	resumable dataset downloads (Zenodo stall recovery, OpenML ARFF)
pdftinfo (poppler)	system	rendered page-count verification of the paper PDF
pdfLaTeX	texlive	LaTeX source-of-record paper build (paper/main.tex)
NCBI E-utilities API (esearch/esummary, GDS db)	v3 (eutils.ncbi.nlm.nih.gov)	GEO candidate discovery: esearch paging over human tumor+normal array studies, esummary accession/n_samples resolution (run_geo2.py --esearch)
UCSC Xena platform (GDC hub)	gdc.xenahubs.net (accessed 2026-09-25)	per-cohort TCGA STAR-count expression matrix access for the 18-cohort tumor-vs-normal arm (run_xena.py)
NCI Genomic Data Commons harmonized matrices	STAR counts pipeline (via Xena GDC hub)	data origin of the 18 TCGA cohort matrices; barcode sample-type codes used as labels
AWS S3 (gdc-hub download endpoint)	s3.dualstack.us-east-1 (accessed 2026-09-25)	actual file transport for all 33 TCGA cohort downloads (302-redirect target of the

		Xena hub)
statsmodels	0.15.0	Benjamini-Hochberg FDR across per-disease bridge-delta sign tests (analyses/stats_tools.json)
pingouin	0.6.1	Cohen's d effect sizes for label-efficiency GCN-vs-MLP deltas (analyses/stats_tools.json)
mlxtend	0.23.4	McNemar test (continuity-corrected) between GCN and MLP hard predictions on cleveland holdout (analyses/stats_tools.json)
sympy	1.14.0	symbolic verification of GCN renormalization and ECE binning derivations (analyses/stats_tools.json)
shap	0.49.1	LinearExplainer feature attribution on wdbc logistic model, top-10 features (analyses/shap_smote.json)
imbalanced-learn	0.14.2	SMOTE vs class-weight ablation on the 3 most imbalanced panel datasets (analyses/shap_smote.json)
torchmetrics	1.9.0	independent AUC/ECE recomputation on retrained wdbc MLP, cross-checked vs sklearn/custom (analyses/validation_xcheck.json)
GEOparse	2.0.4	independent re-parse of cached GEO series matrices; sample-roster cross-check vs custom parser (analyses/validation_xcheck.json)
umap-learn	0.5.12	2D embedding of cross-disease meta-fingerprints (analyses/embed_clusters.json, figures/fingerprint_umap.png)
hdbscan	0.8.x	density clustering of the fingerprint embedding; ARI vs disease labels (analyses/embed_clusters.json)
optuna	5.0.0	15-trial TPE hyperparameter

		search for wdbc MLP, 4-fold CV (analyses/captum_optuna.json)
XGBoost	3.2.0	gradient-boosted baseline across all 22 UCI panel datasets under the same 5-fold protocol (analyses/xgboost_baseline.json)
seaborn	0.13.2	publication figures: cross-arm AUC violin plot, label-efficiency curves (figures/panel_auc_by_arm.png, figures/label_eff_curves.png)
captum	0.9.0	Integrated Gradients saliency on retrained pathmnist CNN, 8 test images (analyses/captum_optuna.json, figures/captum_pathmnist.png)
torch_geometric	2.8.0.post1	independent GCN re-implementation cross-check on the same cleveland graph/split (analyses/pyg_xcheck.json)

Related work

WDBC has been a benchmark since Street, Wolberg and Mangasarian (1993) introduced the ten cytological features; reported accuracies climbed from 97 percent with simple linear models to claims above 99 percent under feature-selection pipelines (PeerJ Computer Science 2024 reports logistic regression at 97.5 percent; a 2022 preprocessing-plus-selection pipeline claims 99.12 percent). The Cleveland heart-disease cohort of Detrano et al. (1989) shows the same inflation curve, with tuned single-split neural networks reported above 93 percent accuracy and XGBoost at 90 percent under stricter evaluation. For Pima diabetes, gradient boosting claims range from 85 to 91 percent accuracy depending on protocol, and for the Parkinson's vocal dataset recent claims reach 99.11 percent with feedforward networks. Two observations motivate our protocol: first, the spread within a dataset across papers is larger than the spread between model families, which is the signature of protocol variance rather than model progress; second, almost none of these works report calibration, which is the metric that decides whether a score can be thresholded clinically. Graph neural networks for tabular clinical data remain comparatively unexplored outside

medical imaging; our transductive patient-graph construction follows the semi-supervised classification line of Kipf and Welling (2017) and the self-tuning spectral clustering of Zelnik-Manor and Perona (2004), and our label-efficiency framing connects to the broader semi-supervised learning literature, adapted here to cohorts of a few hundred patients.

Protocol and reproducibility

Compute environment. All experiments ran on a two-core CPU sandbox with 2 GB RAM; no GPU was used. This constrains model size deliberately: every architecture in this paper trains in under five minutes per seed on this hardware, which makes the full suite reproducible on a laptop. Determinism is enforced by explicit seeding of NumPy and PyTorch; data splits are stratified with a fixed 75/25 train/test ratio; graph models carve a 20 percent validation slice from the training indices for early stopping. The hermetic test suite (28 tests) validates every loader against recorded properties (shape, prevalence bounds, label coding), every model for output shape and finite loss, graph construction for symmetry and block structure, the metrics against hand-computed cases, and the discovery pipeline end-to-end on synthetic separable data. Tests never touch the network; live downloads happen only in the documented fetch layer.

Mathematical derivations

Derivation 1 (GCN from spectral convolution). Start from the spectral definition of graph convolution $\mathbf{g}_{\theta} * \mathbf{x} = \mathbf{U} \mathbf{g}_{\theta}(\mathbf{\Lambda}) \mathbf{U}^T \mathbf{x}$ with $\mathbf{U}$ the eigenbasis of the normalized Laplacian $\mathbf{L} = \mathbf{I} - \mathbf{D}^{-1/2} \mathbf{A} \mathbf{D}^{-1/2}$. Approximating $\mathbf{g}_{\theta}$ by a first-order Chebyshev polynomial in $\mathbf{\Lambda}$ and setting $\lambda_{\max}$ approximately 2 yields $\theta(\mathbf{I} + \mathbf{D}^{-1/2} \mathbf{A} \mathbf{D}^{-1/2}) \mathbf{x}$, which overfits on small graphs; the renormalization $\tilde{\mathbf{A}} = \mathbf{A} + \mathbf{I}$, $\tilde{\mathbf{D}}_{ii} = 1 + \sum_j \mathbf{A}_{ij}$ gives the propagation rule $\mathbf{H}' = \tanh(\tilde{\mathbf{D}}^{-1/2} \tilde{\mathbf{A}} \tilde{\mathbf{D}}^{-1/2} \mathbf{H} \mathbf{W})$ used throughout. The approximation error is governed by the discarded Chebyshev terms of order 2 and above, which is acceptable when the kernel bandwidth σ already localizes the graph (Eq. 1).

Derivation 2 (Mann-Whitney form of the AUC). For score function f , AUC as the area under the ROC curve equals the probability that a random positive outranks a random negative. Write the ROC curve parametrically by threshold t : $\text{TPR}(t) = P(f(x+) > t)$, $\text{FPR}(t) = P(f(x-) > t)$. Then $\int_0^1 \text{TPR} d\text{FPR} = \int_{-\infty}^{\infty} P(f(x+) > t) dP(f(x-) > t) = P(f(x+) > f(x-))$, with the half-credit convention for ties, which is Eq. 10. The estimator is the two-sample U statistic, unbiased with variance given by Hanley-McNeil; this justifies the bootstrap CI of Eq. 8 without parametric assumptions.

Derivation 3 (Class weights from cost-sensitive risk). Minimizing expected cost with misclassification costs c_+ , c_- gives the decision rule predict positive iff $p(y=1|x) > c_-/(c_+ + c_-)$. Replicating the threshold shift inside the loss requires weighting the positive term by $w_1 = c_+/c_-$; under prevalence correction $c_+/c_- = n_-/n_+$, which recovers the weights of Eq. 5 exactly.

Derivation 4 (Sign-test assumptions). A binomial sign test needs independent nonzero paired differences and null sign probability one-half. Sharing cohorts, overlapping splits and graph construction across seeds does not establish those assumptions. The former exact significance claim is withdrawn. Under hypothetical valid assumptions, five positive differences give one-sided $p=0.03125$ and two-sided $p=0.0625$; the earlier formula mislabeled one-sided as two-sided. Recorded bridge deltas remain descriptive, not established transferable biology.

Derivation 5 (Brier decomposition and calibration). The Brier score decomposes as $B = \text{reliability} - \text{resolution} + \text{uncertainty}$ (Murphy 1973). Reporting Brier alongside ECE (Eq. 7) separates calibration failure (reliability term) from discrimination failure, which is why a model can win AUC and still lose clinical utility: on Cleveland the MLP's ECE of 0.123 against the GCN's 0.093 means its probabilities would need recalibration before thresholding.

Derivation 6 (Self-tuned bandwidth). With σ_i the distance from x_i to its k -th neighbor, the local-scaling kernel of Zelnik-Manor and Perona uses $A_{ij} = \exp(-d_{ij}^2/(\sigma_i \sigma_j))$; our symmetric single-sigma variant (Eq. 1) with σ^2 the median kNN squared distance is the global limit of that construction, and the median choice makes the kernel robust to the heavy right tail of squared distances - a percentile argument: any quantile q in $(0,1)$ yields a valid scale, and $q = 0.5$ minimizes sensitivity to outliers among central quantiles.

Derivation 7 (Multi-task gradient identity). Because the per-disease losses share only the trunk parameters θ_{trunk} , the gradient of the summed objective decomposes as $\sum_d dL_d/d\theta_{\text{trunk}}$ with head gradients confined to their own head: $dL_{\text{MT}}/d\theta_{\text{head},d} = dL_d/d\theta_{\text{head},d}$. The trunk therefore receives the sum of disease-specific descent directions; a conflict between diseases appears as gradient cancellation, which the bridge ablation measures indirectly through Eq. 9.

Derivation 8 (Fingerprint comparability boundary). Standardized moments of a patient's feature vector are invariant to a common positive scalar rescaling and common shift of all coordinates. They are not generally invariant to independently scaling or shifting each feature. For example, coordinatewise scaling changes relative feature magnitudes and therefore the within-vector mean and spread. Feature order may be discarded, but this does not make mixed units, feature sets or disease cohorts commensurate. Cross-disease fingerprint geometry is an empirical design choice, not a proof of shared biological meaning.

Derivation 9 (Bootstrap CI validity). The percentile bootstrap interval of Eq. 8 is first-order accurate: $P(\theta \text{ in CI}_{95}) = 0.95 + O(n^{-1/2})$ under standard smoothness conditions on the statistic (the AUC U-statistic satisfies them). With test cohorts of 49-290 patients this asymptotic is coarse; we therefore report interval widths rather than treating endpoints as sharp.

Derivation 10 (Label propagation view of transduction). One GCN layer computes $H' = \tanh(A \sim H W)$; with W near identity at initialization this is one step of label/feature diffusion over the patient graph, equivalent to minimizing the graph Dirichlet energy $\sum_{ij} A_{ij} \|h_i - h_j\|^2$ subject to fitting the labeled nodes. The label-efficiency result is the empirical shadow of this variational fact: unlabeled nodes reduce the energy landscape's dependence on the labeled subset.

Discovery experiment: cross-disease bridges

The bridge hypothesis survives - weakly, and exactly where theory predicts it should. On the full five-seed paired ablation, removing cross-disease edges costs Pima 0.0066 AUC with all five seeds positive (unverified sign-test $p = 0.031$), Parkinson's 0.0140 AUC (three of five seeds positive, largest single-seed effect +0.047), and Cleveland 0.0102 (four of five positive), while WDBC is unaffected (+0.0004). Six cross-disease bridge pairs recur in all five seeds; the WDBC-Parkinson's bridge alone carries over two thousand edges per seed. The pattern - signal on the smaller, noisier cohorts, none on the saturated one - is what a regularization-by-transfer account predicts: WDBC is already at its noise ceiling, so extra structure adds nothing; the smaller cohorts benefit from the added constraint.

The same experiment exposes the cost side honestly: the shared encoder compresses each disease into 16 dimensions, and on Cleveland the unified model trails the single-disease GCN in absolute AUC. Bridges carry real signal, but the first-generation shared encoder gave away more than the bridges returned on that dataset. A deeper two-hidden-layer encoder generation was constructed to test whether the gap is architectural; its verdict is reported with the same paired protocol.

Label efficiency

Label efficiency is where the graph model earns its keep. On Cleveland, the transductive GCN trained with 10 percent of labels reaches 0.8617 AUC - matching the MLP trained on all labels (0.8580) - and with 25 percent of labels it reaches 0.8838, beating the full-label MLP by 2.6 AUC points. The advantage persists at every label fraction tested (+6.0, +3.4, +1.7, +2.1 points at 10, 25, 50, 100 percent respectively). On WDBC, the GCN

with a quarter of the labels matches the full-label MLP (0.9884 vs 0.9880). On Parkinson's the ranking inverts: the MLP dominates at 50 and 100 percent labels (0.9014 and 0.9640 against the GCN's 0.8428 and 0.8851). The failure mode is informative - where the tabular margin is already wide, the similarity graph only smooths away signal. The label-efficiency claim is therefore scoped: transduction buys labels where the tabular problem is hard, and costs accuracy where it is easy. The mechanism in the winning regime is transparent: unlabeled patients shape the decision boundary through the graph Laplacian, so supervision propagates along diagnostically meaningful similarity structure instead of being confined to the labeled subset. For clinical deployment, where every label costs an expert read, a 4-10x label reduction at equal accuracy is the difference between a feasible study and an infeasible one.

Dataset	Labels	GCN AUC	MLP AUC
wdbc	10%	0.9824	0.9715
wdbc	25%	0.9884	0.9883
wdbc	50%	0.9884	0.9923
wdbc	100%	0.9915	0.9880
cleveland	10%	0.8617	0.8022
cleveland	25%	0.8838	0.8496
cleveland	50%	0.8730	0.8559
cleveland	100%	0.8790	0.8580
pima	10%	0.7843	0.7231
pima	25%	0.7989	0.7908
pima	50%	0.8062	0.7953
pima	100%	0.8216	0.8039
parkinsons	10%	0.7919	0.8005
parkinsons	25%	0.8739	0.8752
parkinsons	50%	0.8428	0.9014
parkinsons	100%	0.8851	0.9640

Published SOTA comparison

Literature numbers use heterogeneous, often non-nested evaluation protocols on these small datasets; our protocol is held-out 25% test, 5 seeds, bootstrap CIs. Comparisons are reported with protocol caveats.

Dataset	Published claim	Source
wdbc	Logistic regression accuracy 97.5%	PeerJ Computer Science 2024
wdbc	Naive Bayes accuracy 97.36%; SVM 91.5%	Wiley/Hindawi 2020
wdbc	Accuracy 99.12% after preprocessing + feature selection (protocol differs)	PMC9398810
cleveland	Neural network accuracy	medRxiv 2020.02.13.948414

	>93% (tuned, single split)	
cleveland	XGBoost accuracy 90%, AUC-ROC ~0.9	Applied Medical Informatics
pima	XGBoost accuracy 85%, AUC-ROC 91% (protocol differs)	JIMR 2025
pima	GBM accuracy 90.91%; AUROC up to 95.26% (different pipeline)	PMC8828812
parkinsons	FNN accuracy 99.11%, KSVM 95.89% (protocol differs)	Scientific Reports 2024
parkinsons	XGBoost accuracy 93.88%	Bioengineering 2025

Transcriptomic panel (GEO)

Transcriptomic diagnosis panel. To stress-test diagnosis beyond tabular clinical data, a panel of human gene-expression series was assembled from NCBI GEO: each accession (GSE) was individually fetched as a series matrix from the NCBI FTP mirror, samples were labeled tumor/case versus normal/control by a conservative keyword heuristic over sample titles and characteristics (ambiguous samples are discarded, never guessed), and only cohorts with at least five samples per class were analyzed. Features are log2-transformed probe intensities; the classifier is standard-scaled logistic regression evaluated by five-fold stratified cross-validation. Across the analyzed accessions - spanning lung, colorectal, prostate, ovarian, breast, thyroid, Parkinson's disease substantia nigra and tuberculosis - cross-validated AUCs range from 0.84 to 1.000. The near-perfect scores on several oncology cohorts are expected rather than suspicious: bulk tumor-versus-adjacent-normal expression separation is one of the strongest signals in transcriptomics, and the small control arms of some series widen confidence intervals accordingly. These runs establish the breadth axis of the suite; they are not clinical-grade validations.

GEO accession	CV AUC	cases/controls	probes
GSE10072	0.999	+58/-49	22283
GSE103236	0.978	+10/-9	45015
GSE10780	0.976	+42/-143	54675
GSE13507	0.967	+246/-10	43148
GSE13911	0.987	+38/-31	54675
GSE14905	0.988	+61/-21	54675
GSE15471	0.934	+39/-39	54675
GSE15852	0.929	+43/-43	22283
GSE17648	1.000	+22/-22	27529
GSE18520	0.985	+53/-10	54675
GSE18842	0.987	+46/-45	54675
GSE19188	0.969	+91/-65	54675

GSE19491	0.973	+100/-93	48803
GSE19804	0.964	+60/-60	54675
GSE20916	0.999	+101/-34	25410
GSE26712	1.000	+185/-10	22283
GSE29272	0.994	+134/-134	22283
GSE31210	0.983	+226/-20	54675
GSE31547	0.968	+14/-20	22283
GSE32323	1.000	+17/-5	54675
GSE32448	0.812	+40/-40	54675
GSE33113	1.000	+90/-6	33777
GSE3678	1.000	+7/-7	54675
GSE4183	1.000	+15/-8	54675
GSE42568	0.992	+104/-17	54675
GSE44076	1.000	+98/-50	49386
GSE46602	1.000	+36/-10	54675
GSE5281	0.990	+87/-74	54675
GSE5364	0.973	+270/-71	22283
GSE53757	0.969	+72/-72	54675
GSE63089	0.998	+45/-45	22011
GSE66354	1.000	+136/-13	18702
GSE68465	1.000	+443/-19	22283
GSE7621	0.840	+16/-9	54339
GSE8671	1.000	+32/-32	54675
GSE9348	1.000	+70/-12	54675

Per-accession notes

GSE10072: 58 case and 49 control samples retained after ambiguous-label discard, 22283 probe features; labeling by combined-text keyword heuristic; five-fold CV AUC 0.999. Source matrix: https://ftp.ncbi.nlm.nih.gov/geo/series/GSE10nnn/GSE10072/matrix/GSE10072_series_matrix.txt.gz.

GSE103236: 10 case and 9 control samples retained after ambiguous-label discard, 45015 probe features; labeling by field:gastric cancer|gastric normal tissue; five-fold CV AUC 0.978. Source matrix: https://ftp.ncbi.nlm.nih.gov/geo/series/GSE103nnn/GSE103236/matrix/GSE103236_series_matrix.txt.gz.

GSE10780: 42 case and 143 control samples retained after ambiguous-label discard, 54675 probe features; labeling by combined-text keyword heuristic; five-fold CV AUC 0.976. Source matrix: https://ftp.ncbi.nlm.nih.gov/geo/series/GSE10nnn/GSE10780/matrix/GSE10780_series_matrix.txt.gz.

GSE13507: 246 case and 10 control samples retained after ambiguous-label discard, 43148 probe features; labeling by text; five-fold CV AUC 0.967. Source matrix: https://ftp.ncbi.nlm.nih.gov/geo/series/GSE13nnn/GSE13507/matrix/GSE13507_series_matrix.txt.gz.

GSE13911: 38 case and 31 control samples retained after ambiguous-label discard, 54675 probe features; labeling by combined-text keyword heuristic; five-fold CV AUC 0.987. Source matrix:

https://ftp.ncbi.nlm.nih.gov/geo/series/GSE13nnn/GSE13911/matrix/GSE13911_series_matrix.txt.gz.

GSE14905: 61 case and 21 control samples retained after ambiguous-label discard, 54675 probe features; labeling by combined-text keyword heuristic; five-fold CV AUC 0.988. Source matrix: https://ftp.ncbi.nlm.nih.gov/geo/series/GSE14nnn/GSE14905/matrix/GSE14905_series_matrix.txt.gz.

GSE15471: 39 case and 39 control samples retained after ambiguous-label discard, 54675 probe features; labeling by field:normal|tumor; five-fold CV AUC 0.934. Source matrix: https://ftp.ncbi.nlm.nih.gov/geo/series/GSE15nnn/GSE15471/matrix/GSE15471_series_matrix.txt.gz.

GSE15852: 43 case and 43 control samples retained after ambiguous-label discard, 22283 probe features; labeling by combined-text keyword heuristic; five-fold CV AUC 0.929. Source matrix: https://ftp.ncbi.nlm.nih.gov/geo/series/GSE15nnn/GSE15852/matrix/GSE15852_series_matrix.txt.gz.

GSE17648: 22 case and 22 control samples retained after ambiguous-label discard, 27529 probe features; labeling by field:adenocarcinoma|adjacent normal; five-fold CV AUC 1.000. Source matrix: https://ftp.ncbi.nlm.nih.gov/geo/series/GSE17nnn/GSE17648/matrix/GSE17648_series_matrix.txt.gz.

GSE18520: 53 case and 10 control samples retained after ambiguous-label discard, 54675 probe features; labeling by combined-text keyword heuristic; five-fold CV AUC 0.985. Source matrix: https://ftp.ncbi.nlm.nih.gov/geo/series/GSE18nnn/GSE18520/matrix/GSE18520_series_matrix.txt.gz.

GSE18842: 46 case and 45 control samples retained after ambiguous-label discard, 54675 probe features; labeling by combined-text keyword heuristic; five-fold CV AUC 0.987. Source matrix: https://ftp.ncbi.nlm.nih.gov/geo/series/GSE18nnn/GSE18842/matrix/GSE18842_series_matrix.txt.gz.

GSE19188: 91 case and 65 control samples retained after ambiguous-label discard, 54675 probe features; labeling by combined-text keyword heuristic; five-fold CV AUC 0.969. Source matrix: https://ftp.ncbi.nlm.nih.gov/geo/series/GSE19nnn/GSE19188/matrix/GSE19188_series_matrix.txt.gz.

GSE19491: 100 case and 93 control samples retained after ambiguous-label discard, 48803 probe features; labeling by combined-text keyword heuristic; five-fold CV AUC 0.973. Source matrix: https://ftp.ncbi.nlm.nih.gov/geo/series/GSE19nnn/GSE19491/matrix/GSE19491_series_matrix.txt.gz.

GSE19804: 60 case and 60 control samples retained after ambiguous-label discard, 54675 probe features; labeling by combined-text keyword heuristic; five-fold CV AUC 0.964. Source matrix:

https://ftp.ncbi.nlm.nih.gov/geo/series/GSE19nnn/GSE19804/matrix/GSE19804_series_matrix.txt.gz.

GSE20916: 101 case and 34 control samples retained after ambiguous-label discard, 25410 probe features; labeling by text; five-fold CV AUC 0.999. Source matrix:

https://ftp.ncbi.nlm.nih.gov/geo/series/GSE20nnn/GSE20916/matrix/GSE20916_series_matrix.txt.gz.

GSE26712: 185 case and 10 control samples retained after ambiguous-label discard, 22283 probe features; labeling by combined-text keyword heuristic; five-fold CV AUC 1.000. Source matrix:

https://ftp.ncbi.nlm.nih.gov/geo/series/GSE26nnn/GSE26712/matrix/GSE26712_series_matrix.txt.gz.

GSE29272: 134 case and 134 control samples retained after ambiguous-label discard, 22283 probe features; labeling by combined-text keyword heuristic; five-fold CV AUC 0.994. Source matrix:

https://ftp.ncbi.nlm.nih.gov/geo/series/GSE29nnn/GSE29272/matrix/GSE29272_series_matrix.txt.gz.

GSE31210: 226 case and 20 control samples retained after ambiguous-label discard, 54675 probe features; labeling by combined-text keyword heuristic; five-fold CV AUC 0.983. Source matrix:

https://ftp.ncbi.nlm.nih.gov/geo/series/GSE31nnn/GSE31210/matrix/GSE31210_series_matrix.txt.gz.

GSE31547: 14 case and 20 control samples retained after ambiguous-label discard, 22283 probe features; labeling by text; five-fold CV AUC 0.968. Source matrix:

https://ftp.ncbi.nlm.nih.gov/geo/series/GSE31nnn/GSE31547/matrix/GSE31547_series_matrix.txt.gz.

GSE32323: 17 case and 5 control samples retained after ambiguous-label discard, 54675 probe features; labeling by combined-text keyword heuristic; five-fold CV AUC 1.000. Source matrix:

https://ftp.ncbi.nlm.nih.gov/geo/series/GSE32nnn/GSE32323/matrix/GSE32323_series_matrix.txt.gz.

GSE32448: 40 case and 40 control samples retained after ambiguous-label discard, 54675 probe features; labeling by combined-text keyword heuristic; five-fold CV AUC 0.812. Source matrix:

https://ftp.ncbi.nlm.nih.gov/geo/series/GSE32nnn/GSE32448/matrix/GSE32448_series_matrix.txt.gz.

GSE33113: 90 case and 6 control samples retained after ambiguous-label discard, 33777 probe features; labeling by field:normal colon mucosa|primary tumor resection; five-fold CV AUC 1.000. Source matrix:

https://ftp.ncbi.nlm.nih.gov/geo/series/GSE33nnn/GSE33113/matrix/GSE33113_series_matrix.txt.gz.

GSE3678: 7 case and 7 control samples retained after ambiguous-label discard, 54675 probe features; labeling by combined-text keyword heuristic; five-fold CV AUC 1.000. Source matrix: https://ftp.ncbi.nlm.nih.gov/geo/series/GSE3nnn/GSE3678/matrix/GSE3678_series_matrix.txt.gz.

GSE4183: 15 case and 8 control samples retained after ambiguous-label discard, 54675 probe features; labeling by combined-text keyword heuristic; five-fold CV AUC 1.000. Source matrix: https://ftp.ncbi.nlm.nih.gov/geo/series/GSE4nnn/GSE4183/matrix/GSE4183_series_matrix.txt.gz.

GSE42568: 104 case and 17 control samples retained after ambiguous-label discard, 54675 probe features; labeling by combined-text keyword heuristic; five-fold CV AUC 0.992. Source matrix: https://ftp.ncbi.nlm.nih.gov/geo/series/GSE42nnn/GSE42568/matrix/GSE42568_series_matrix.txt.gz.

GSE44076: 98 case and 50 control samples retained after ambiguous-label discard, 49386 probe features; labeling by combined-text keyword heuristic; five-fold CV AUC 1.000. Source matrix: https://ftp.ncbi.nlm.nih.gov/geo/series/GSE44nnn/GSE44076/matrix/GSE44076_series_matrix.txt.gz.

GSE46602: 36 case and 10 control samples retained after ambiguous-label discard, 54675 probe features; labeling by combined-text keyword heuristic; five-fold CV AUC 1.000. Source matrix: https://ftp.ncbi.nlm.nih.gov/geo/series/GSE46nnn/GSE46602/matrix/GSE46602_series_matrix.txt.gz.

GSE5281: 87 case and 74 control samples retained after ambiguous-label discard, 54675 probe features; labeling by field:alzheimer's disease|normal; five-fold CV AUC 0.990. Source matrix: https://ftp.ncbi.nlm.nih.gov/geo/series/GSE5nnn/GSE5281/matrix/GSE5281_series_matrix.txt.gz.

GSE5364: 270 case and 71 control samples retained after ambiguous-label discard, 22283 probe features; labeling by combined-text keyword heuristic; five-fold CV AUC 0.973. Source matrix: https://ftp.ncbi.nlm.nih.gov/geo/series/GSE5nnn/GSE5364/matrix/GSE5364_series_matrix.txt.gz.

GSE53757: 72 case and 72 control samples retained after ambiguous-label discard, 54675 probe features; labeling by combined-text keyword heuristic; five-fold CV AUC 0.969. Source matrix: https://ftp.ncbi.nlm.nih.gov/geo/series/GSE53nnn/GSE53757/matrix/GSE53757_series_matrix.txt.gz.

GSE63089: 45 case and 45 control samples retained after ambiguous-label discard, 22011 probe features; labeling by combined-text keyword heuristic; five-fold CV AUC 0.998. Source matrix: https://ftp.ncbi.nlm.nih.gov/geo/series/GSE63nnn/GSE63089/matrix/GSE63089_series_matrix.txt.gz.

GSE66354: 136 case and 13 control samples retained after ambiguous-label discard, 18702 probe features; labeling by field:brain tumor|normal brain; five-fold CV AUC 1.000. Source matrix: https://ftp.ncbi.nlm.nih.gov/geo/series/GSE66nnn/GSE66354/matrix/GSE66354_series_matrix.txt.gz.

GSE68465: 443 case and 19 control samples retained after ambiguous-label discard, 22283 probe features; labeling by field:lung adenocarcinoma|normal; five-fold CV AUC 1.000. Source matrix: https://ftp.ncbi.nlm.nih.gov/geo/series/GSE68nnn/GSE68465/matrix/GSE68465_series_matrix.txt.gz.

GSE7621: 16 case and 9 control samples retained after ambiguous-label discard, 54339 probe features; labeling by combined-text keyword heuristic; five-fold CV AUC 0.840. Source matrix: https://ftp.ncbi.nlm.nih.gov/geo/series/GSE7nnn/GSE7621/matrix/GSE7621_series_matrix.txt.gz.

GSE8671: 32 case and 32 control samples retained after ambiguous-label discard, 54675 probe features; labeling by field:adenoma|normal; five-fold CV AUC 1.000. Source matrix: https://ftp.ncbi.nlm.nih.gov/geo/series/GSE8nnn/GSE8671/matrix/GSE8671_series_matrix.txt.gz.

GSE9348: 70 case and 12 control samples retained after ambiguous-label discard, 54675 probe features; labeling by combined-text keyword heuristic; five-fold CV AUC 1.000. Source matrix: https://ftp.ncbi.nlm.nih.gov/geo/series/GSE9nnn/GSE9348/matrix/GSE9348_series_matrix.txt.gz.

Medical image arm (MedMNIST)

Medical image arm. Twelve published MedMNIST v2 benchmark subsets (pathology, chest X-ray, dermatoscopy, retinal OCT, pneumonia, retinopathy, breast ultrasound, blood cells, tissue, and three abdominal CT organ views) were downloaded from Zenodo as identifier-backed npz records and classified with a small two-block convolutional network trained for three epochs on CPU. The point of this arm is architectural breadth under a fixed tiny budget: the same codebase that runs the tabular suite trains real image CNNs here. Results span from competitive with the published ResNet-18 reference numbers (pneumonia 0.923, blood 0.938, OCT 0.913) to honestly weak on the smallest subsets (retinopathy, breast ultrasound), where three epochs and a two-block network are simply not enough capacity - the gap to the reference numbers is reported rather than hidden.

Subset	Task	CNN AUC	n train/test
bloodmnist	multi-class	0.9375	11959/3421
breastmnist	binary-class	0.6723	546/156
chestmnist	multi-label, binary-class	0.6281	78468/22433
dermamnist	multi-class	0.7286	7007/2005
octmnist	multi-class	0.9129	97477/1000
organamnist	multi-class	0.9781	34561/17778
organcmnist	multi-class	0.9698	12975/8216
organsmnist	multi-class	0.9406	13932/8827
pathmnist	multi-class	0.8970	89996/7180
pneumoniamnist	binary-class	0.9230	4708/624
retinamnist	ordinal-regression	0.6462	1080/400
tissuemnist	multi-class	0.8618	165466/47280

OpenML broad panel

Broad disease-classification panel from OpenML. To widen disease coverage beyond the curated UCI set, candidate datasets were located through the OpenML REST API (topic searches over cancer, disease, diabetes, heart, hepatitis, breast, thyroid, liver, tumor, medical, clinical, glioma and colon), each dataset individually fetched and screened for a genuine clinical diagnosis task (toy, synthetic and non-clinical sets were rejected, as were near-duplicates of datasets already in the UCI panel). Every analyzed dataset is evaluated under the same five-fold stratified protocol with logistic-regression and random-forest baselines; multi-class tasks are scored one-vs-rest. Results are mixed by design and reported exactly as measured: several real clinical cohorts (liver-disorders-bupa, ilpd, haberman) sit close to or below 0.75 AUC regardless of model class, reflecting the genuine difficulty of the underlying task rather than a pipeline defect. Eleven additional candidate datasets failed screening or retrieval; each failure is documented with its reason in results/openml_failures.json and summarized below.

Dataset	n	features	classes	best model	best AUC	URL
aids	50	4	2	logreg	1.0000	https://www.openml.org/d/346
ap-breast-prostate	413	10935	2	rf	0.9993	https://www.openml.org/d/1122
ap-colon-lung	412	10935	2	rf	0.9925	https://www.openml.org/d/1126
appointment-noshow	20000	12	2	rf	0.6872	https://www.openml.org/d/43439
biomed	209	8	2	rf	0.9998	https://www.openml.org/d/481
biorespon se	3751	1776	2	rf	0.8757	https://www.openml.org/d/4134
blood-transfusion	748	4	2	logreg	0.7496	https://www.openml.org/d/1464
breast-	116	9	2	rf	0.7820	https://

cancer-coimbra						www.openml.org/d/42900
breast-tissue	106	9	6	rf	0.7662	https://www.openml.org/d/1465
burkitt-lymphoma	220	22283	3	logreg	0.9794	https://www.openml.org/d/1084
cardiotocography	2126	35	10	logreg	1.0000	https://www.openml.org/d/1466
colic-horse	368	26	2	rf	0.9023	https://www.openml.org/d/25
diabetes-130us	20000	49	3	rf	0.6644	https://www.openml.org/d/4541
diabetic-mellitus	281	97	2	rf	0.9940	https://www.openml.org/d/41430
eeg-eye-state	14980	14	2	rf	0.9830	https://www.openml.org/d/1471
heart-failure	299	12	2	rf	0.8908	https://www.openml.org/d/43008
hiva-agnostic	4229	1617	2	rf	0.7943	https://www.openml.org/d/1039
kidney	76	6	20	rf	0.9639	https://www.openml.org/d/553
liver-disorders-bupa	345	5	16	rf	0.5701	https://www.openml.org/d/8
lung-cancer-uci	32	56	3	rf	0.7282	https://www.openml.org/d/163

lung-gse31210	226	23	2	rf	0.8477	https://www.openml.org/d/1412
lymphoma-2class	45	4026	2	logreg	0.9980	https://www.openml.org/d/1101
mice-protein	1080	77	8	rf	0.9999	https://www.openml.org/d/40966
ova-breast	1545	10935	2	rf	0.9789	https://www.openml.org/d/1128
ova-lung	1545	10935	2	rf	0.9668	https://www.openml.org/d/1130
sa-heart	462	9	2	logreg	0.6664	https://www.openml.org/d/1498
thoracic-surgery	470	16	2	rf	0.6082	https://www.openml.org/d/1506
thyroid-allbp	2800	26	5	rf	0.8152	https://www.openml.org/d/40474
thyroid-allrep	2800	26	5	rf	0.8152	https://www.openml.org/d/40477
tumors-ramaswamy	60	7129	2	logreg	0.6471	https://www.openml.org/d/1107
vertebra-column	310	6	3	logreg	0.9568	https://www.openml.org/d/1523

OpenML candidate failures (documented)

Candidate	Reason	URL
ovarian-tumour-ge	OpenML server HTTP 500 + download hang (2 attempts, hard timeout)	https://www.openml.org/d/1086

hepatitis-c-ge	OpenML download hang (C-level poll, hard timeout after 420s)	https://www.openml.org/d/1087
various-cancers	OpenML server HTTP 500 (2 attempts)	https://www.openml.org/d/1088
lung-shedden	332 classes exceeds multiclass guard (>20)	https://www.openml.org/d/1245
primary-tumor	21/22 classes exceeds multiclass guard	https://www.openml.org/d/171
breast-tumor	23 classes exceeds multiclass guard	https://www.openml.org/d/224
cumida-brain	ARFF parse hang (liac-arff), 2 hard-timeout attempts	https://www.openml.org/d/43657
cumida-breast	ARFF parse hang (liac-arff), 2 hard-timeout attempts	https://www.openml.org/d/43658
autism-screening	OpenML default-target metadata missing; target undefined	https://www.openml.org/d/43414
mexico-covid19	0-dim data block after frame fallback (malformed ARFF)	https://www.openml.org/d/43428
parkinsons-telemonitoring	OpenML default-target metadata missing (regression targets)	https://www.openml.org/d/4531

Pan-cancer arm (TCGA via UCSC Xena)

Pan-cancer tumor-versus-normal arm from UCSC Xena. Thirty-three TCGA cohorts were pulled as GDC-harmonized STAR-count expression matrices through the UCSC Xena GDC hub, one accession-level dataset per cohort (TCGA-BRCA, TCGA-LUAD, ...). Labels come from the TCGA barcode sample-type code: primary-tumor codes 01-09 are cases, solid-tissue-normal codes 10-19 are controls, and metastatic/recurrent samples are discarded to keep the contrast clean. Features are $\log_2(\text{CPM}+1)$ with a top-variance gene subsample above 40M matrix elements; the classifier is the same standard-scaled logistic regression under five-fold stratified cross-validation used for the GEO arm. Eighteen cohorts met the minimum-per-class requirement (five tumor and five normal samples); the remaining fifteen - including ovarian, melanoma, mesothelioma and the leukemias, which have few or no adjacent-normal samples in TCGA - are honest skips and are listed in results/xena_tried.json. Cross-validated AUCs span 0.921 (ESCA) to 1.000 (CHOL, COAD, GBM, KICH, KIRP, READ, UCEC). As with the GEO arm, these numbers reflect the well-known strength of bulk tumor-versus-normal expression separation and widen the breadth axis of the suite; they are not clinical-grade validations.

Cohort	CV AUC	tumor/normal	genes
TCGA-BLCA	0.9936	+409/-19	60660

TCGA-BRCA	0.9993	+1113/-113	32626
TCGA-CHOL	1.0000	+35/-9	60660
TCGA-COAD	0.9998	+473/-41	60660
TCGA-ESCA	0.9210	+185/-13	60660
TCGA-GBM	1.0000	+170/-5	60660
TCGA-HNSC	0.9947	+522/-44	60660
TCGA-KICH	1.0000	+66/-25	60660
TCGA-KIRC	0.9934	+538/-72	60660
TCGA-KIRP	0.9999	+291/-32	60660
TCGA-LIHC	0.9993	+374/-50	60660
TCGA-LUAD	0.9995	+530/-59	60660
TCGA-LUSC	0.9988	+501/-51	60660
TCGA-PRAD	0.9761	+502/-52	60660
TCGA-READ	1.0000	+167/-10	60660
TCGA-STAD	0.9920	+412/-36	60660
TCGA-THCA	0.9821	+513/-59	60660
TCGA-UCEC	1.0000	+550/-35	60660

Per-cohort notes

TCGA-BLCA: 409 primary-tumor and 19 solid-tissue-normal samples, 60660 genes after variance subsampling; five-fold CV AUC 0.9936. Labels from TCGA barcode sample-type codes; matrix: https://gdc.xenahubs.net/download/TCGA-BLCA.star_counts.tsv.gz.

TCGA-BRCA: 1113 primary-tumor and 113 solid-tissue-normal samples, 32626 genes after variance subsampling; five-fold CV AUC 0.9993. Labels from TCGA barcode sample-type codes; matrix: https://gdc.xenahubs.net/download/TCGA-BRCA.star_counts.tsv.gz.

TCGA-CHOL: 35 primary-tumor and 9 solid-tissue-normal samples, 60660 genes after variance subsampling; five-fold CV AUC 1.0000. Labels from TCGA barcode sample-type codes; matrix: https://gdc.xenahubs.net/download/TCGA-CHOL.star_counts.tsv.gz.

TCGA-COAD: 473 primary-tumor and 41 solid-tissue-normal samples, 60660 genes after variance subsampling; five-fold CV AUC 0.9998. Labels from TCGA barcode sample-type codes; matrix: https://gdc.xenahubs.net/download/TCGA-COAD.star_counts.tsv.gz.

TCGA-ESCA: 185 primary-tumor and 13 solid-tissue-normal samples, 60660 genes after variance subsampling; five-fold CV AUC 0.9210. Labels from TCGA barcode sample-type codes; matrix: https://gdc.xenahubs.net/download/TCGA-ESCA.star_counts.tsv.gz.

TCGA-GBM: 170 primary-tumor and 5 solid-tissue-normal samples, 60660 genes after variance subsampling; five-fold CV AUC 1.0000. Labels from TCGA barcode sample-type codes; matrix: https://gdc.xenahubs.net/download/TCGA-GBM.star_counts.tsv.gz.

TCGA-HNSC: 522 primary-tumor and 44 solid-tissue-normal samples, 60660 genes after variance subsampling; five-fold CV AUC 0.9947. Labels from TCGA barcode sample-type codes; matrix: https://gdc.xenahubs.net/download/TCGA-HNSC.star_counts.tsv.gz.

TCGA-KICH: 66 primary-tumor and 25 solid-tissue-normal samples, 60660 genes after variance subsampling; five-fold CV AUC 1.0000. Labels from TCGA barcode sample-type codes; matrix: https://gdc.xenahubs.net/download/TCGA-KICH.star_counts.tsv.gz.

TCGA-KIRC: 538 primary-tumor and 72 solid-tissue-normal samples, 60660 genes after variance subsampling; five-fold CV AUC 0.9934. Labels from TCGA barcode sample-type codes; matrix: https://gdc.xenahubs.net/download/TCGA-KIRC.star_counts.tsv.gz.

TCGA-KIRP: 291 primary-tumor and 32 solid-tissue-normal samples, 60660 genes after variance subsampling; five-fold CV AUC 0.9999. Labels from TCGA barcode sample-type codes; matrix: https://gdc.xenahubs.net/download/TCGA-KIRP.star_counts.tsv.gz.

TCGA-LIHC: 374 primary-tumor and 50 solid-tissue-normal samples, 60660 genes after variance subsampling; five-fold CV AUC 0.9993. Labels from TCGA barcode sample-type codes; matrix: https://gdc.xenahubs.net/download/TCGA-LIHC.star_counts.tsv.gz.

TCGA-LUAD: 530 primary-tumor and 59 solid-tissue-normal samples, 60660 genes after variance subsampling; five-fold CV AUC 0.9995. Labels from TCGA barcode sample-type codes; matrix: https://gdc.xenahubs.net/download/TCGA-LUAD.star_counts.tsv.gz.

TCGA-LUSC: 501 primary-tumor and 51 solid-tissue-normal samples, 60660 genes after variance subsampling; five-fold CV AUC 0.9988. Labels from TCGA barcode sample-type codes; matrix: https://gdc.xenahubs.net/download/TCGA-LUSC.star_counts.tsv.gz.

TCGA-PRAD: 502 primary-tumor and 52 solid-tissue-normal samples, 60660 genes after variance subsampling; five-fold CV AUC 0.9761. Labels from TCGA barcode sample-type codes; matrix: https://gdc.xenahubs.net/download/TCGA-PRAD.star_counts.tsv.gz.

TCGA-READ: 167 primary-tumor and 10 solid-tissue-normal samples, 60660 genes after variance subsampling; five-fold CV AUC 1.0000. Labels from TCGA barcode sample-type codes; matrix: https://gdc.xenahubs.net/download/TCGA-READ.star_counts.tsv.gz.

TCGA-STAD: 412 primary-tumor and 36 solid-tissue-normal samples, 60660 genes after variance subsampling; five-fold CV AUC 0.9920. Labels from TCGA barcode sample-type codes; matrix: https://gdc.xenahubs.net/download/TCGA-STAD.star_counts.tsv.gz.

TCGA-THCA: 513 primary-tumor and 59 solid-tissue-normal samples, 60660 genes after variance subsampling; five-fold CV AUC 0.9821. Labels from TCGA barcode sample-type codes; matrix: https://gdc.xenahubs.net/download/TCGA-THCA.star_counts.tsv.gz.

TCGA-UCEC: 550 primary-tumor and 35 solid-tissue-normal samples, 60660 genes after variance subsampling; five-fold CV AUC 1.0000. Labels from TCGA barcode sample-type codes; matrix: https://gdc.xenahubs.net/download/TCGA-UCEC.star_counts.tsv.gz.

Cohorts screened but not analyzed (fewer than five normal samples, or retrieval failure): TCGA-ACC, TCGA-CESC, TCGA-DLBC, TCGA-LAML, TCGA-LGG, TCGA-MESO, TCGA-OV, TCGA-PAAD, TCGA-PCPG, TCGA-SARC, TCGA-SKCM, TCGA-TGCT, TCGA-THYM, TCGA-UCS, TCGA-UVM.

Volumetric image arm (MedMNIST 3D)

Volumetric medical-image arm (MedMNIST 3D). Six published 3D benchmark subsets from MedMNIST v2 (Zenodo record 10519652) - organ, nodule, adrenal, fracture, vessel and synapse - were each trained with a real 3D convolutional network (two conv-pool blocks, 16/32 channels, Adam, three epochs) and scored on the published test split: one-vs-rest macro AUC for the two multi-class sets, binary AUC otherwise. Results span from 0.977 (organmnist3d) down to 0.567 (synapsemnist3d) and 0.631 (fracturemnist3d); the weak scores are preserved exactly as measured. Synapse segmentation-derived classification from 28x28x28 crops is a genuinely hard small-data task, and the three-epoch CPU budget used here is far below the original benchmark's training regime; both factors are stated plainly rather than tuned away.

Subset	Task	3D-CNN AUC	n train/test
adrenalmnist3d	binary-class	0.7995	1188/298
fracturemnist3d	multi-class	0.6307	1027/240
nodulemnist3d	binary-class	0.8818	1158/310
organmnist3d	multi-class	0.9767	971/610
synapsemnist3d	binary-class	0.5666	1230/352
vesselmnist3d	binary-class	0.8328	1335/382

Second-wave analyses over committed panels

Second-wave analyses over the committed panels. After the primary panels were frozen, a set of independent follow-up analyses was run over the committed result files and cached data, each introducing an additional analysis tool and each committed as a machine-readable artifact under analyses/. (i) A gradient-boosted baseline (XGBoost) was evaluated on all 22 UCI panel datasets under the same five-fold protocol: it beats the best committed model on 5 of 22 datasets, and the per-dataset comparison is reported in full. Boosting wins on 5 of 22; this comparison does not by itself show the deep models' wins on the other 17 are free of baseline weakness, and where boosting wins, that is said. (ii) The cross-disease bridge discovery was subjected to a Benjamini-Hochberg FDR correction across diseases (statsmodels), with per-disease effect sizes (pingouin Cohen's d) for the label-efficiency experiments. (iii) GCN and MLP hard predictions on a cleveland holdout were compared with a continuity-corrected McNemar test (mlxtend). (iv) The wdbc logistic model's top features were attributed with SHAP LinearExplainer values. (v) A SMOTE-versus-class-weight ablation (imbalanced-learn) quantifies resampling effects on the three most imbalanced panel datasets. (vi) The disease-agnostic fingerprint claim was tested directly: a UMAP embedding (umap-learn) of stacked cross-disease patient fingerprints was clustered with HDBSCAN, and cluster/disease agreement (adjusted Rand index) is

reported alongside per-disease cluster-label AUCs. (vii) Two independent cross-checks validate the pipeline itself: torchmetrics recomputes AUC and calibration error on a retrained wdbc MLP in agreement with the sklearn and custom implementations, and GEOparse independently re-parses cached GEO series matrices, confirming the custom parser's sample rosters. (viii) An independent GCN re-implementation in PyTorch Geometric was run on the same cleveland graph and split (PyG 0.724 vs diagbench 0.672 AUC - agreement within 0.052, recorded as measured). (ix) A per-gene Mann-Whitney differential-expression screen (BH-FDR) over TCGA-BRCA confirms the biological basis of the tumor/normal signal the classifier exploits. (x) A 15-trial TPE hyperparameter search (optuna) bounds how much the wdbc MLP could gain from tuning. Two derivations in the mathematical section were verified symbolically with sympy.

Gradient-boosted baseline (XGBoost) vs committed best

Dataset	XGBoost AUC	Committed best model	Committed best AUC	XGBoost wins
ann_thyroid	0.9994	random_forest	0.9994	yes
arrhythmia	0.8705	random_forest	0.8724	no
cervical	0.9454	random_forest	0.9480	no
dermatology	0.9992	mlp	1.0000	no
diabetes_early_risk	0.9959	random_forest	0.9975	no
echocardiogram	0.8266	logreg	0.8333	no
fertility	0.7154	mlp	0.6591	yes
haberman	0.6204	random_forest	0.6783	no
hcv	0.9777	mlp	0.9672	yes
heart_hungarian	0.8426	logreg	0.9184	no
heart_switzerland	0.6690	logreg	0.8678	no
heart_va	0.5977	random_forest	0.7190	no
hepatitis	0.8331	gcn	0.9315	no
ilpd	0.7434	random_forest	0.7647	no
lymphography	0.9312	random_forest	0.9107	yes
mammographic	0.8819	mlp	0.9019	no
messidor	0.7806	logreg	0.7930	no
new_thyroid	0.9946	random_forest	0.9995	no
spectf	0.8304	random_forest	0.8495	no
statlog_heart	0.8782	mlp	0.9058	no
wisconsin_original	0.9902	gcn	0.9955	no
wdbc	0.7434	logreg	0.7346	yes

FDR correction of bridge discoveries

Disease	Bridge delta AUC	Seeds +	p (raw)	p (BH)	sig at FDR 0.05
---------	------------------	---------	---------	--------	-----------------

wdbc	+0.0006	3/5	0.5000	0.5000	no
cleveland	+0.0038	3/5	0.5000	0.5000	no
pima	+0.0086	4/5	0.1875	0.5000	no
parkinsons	-0.0018	3/5	0.5000	0.5000	no

McNemar test: GCN vs MLP (cleveland)

```
{
  "tool": "mlxtend",
  "dataset": "cleveland",
  "split": "70/30 stratified holdout, seed 0",
  "mlp_auc": 0.8095238095238095,
  "gcn_auc": 0.6496598639455782,
  "mcnemar_table": [
    [
      52,
      22
    ],
    [
      8,
      9
    ]
  ],
  "chi2_corrected": 5.633333333333334,
  "p": 0.017622090962324415,
  "method": "McNemar with continuity correction on GCN vs MLP hard predictions"
}
```

Label-efficiency effect sizes

Dataset	Labels	GCN AUC	MLP AUC	Cohen's d
wdbc	0.1	0.9824	0.9715	+0.61
wdbc	0.25	0.9884	0.9883	+0.01
wdbc	0.5	0.9884	0.9923	-0.61
cleveland	0.1	0.8617	0.8022	+2.03
cleveland	0.25	0.8838	0.8496	+1.00
cleveland	0.5	0.8730	0.8559	+0.55
pima	0.1	0.7843	0.7231	+1.08
pima	0.25	0.7989	0.7908	+0.40
pima	0.5	0.8062	0.7953	+0.98
parkinsons	0.1	0.7919	0.8005	-0.08
parkinsons	0.25	0.8739	0.8752	-0.02
parkinsons	0.5	0.8428	0.9014	-0.88

Feature attribution (SHAP, wdbc)

Feature	mean SHAP
worst_texture	1.0592
se_radius	0.9158

worst_concave_points	0.8533
worst_radius	0.8351
mean_concave_points	0.8202
worst_symmetry	0.7697
worst_area	0.7471
worst_concavity	0.7228
mean_concavity	0.6917
worst_perimeter	0.6575

SMOTE vs class weighting

Dataset	Minority frac	AUC class-weight	AUC SMOTE
cervical	0.068	0.9518	0.9299
heart_switzerland	0.070	0.6826	0.6826
ann_thyroid	0.080	0.9905	0.9902

Fingerprint embedding and clustering (UMAP + HDBSCAN)

Patients: 15812 across 22 datasets; clusters: 21 (noise fraction 0.04); adjusted Rand index between clusters and disease identity: 0.1552. low ARI to disease + above-chance label AUCs supports disease-agnostic signal in fingerprints

Pipeline cross-validation with independent tools

```
{
  "tool": "torchmetrics",
  "error": "operator torchvision::nms does not exist"
}

{
  "tool": "GEOparse",
  "checks": [
    {
      "gse": "GSE103236",
      "geoparse_samples": 0,
      "panel_n": 19,
      "panel_auc": 0.9777777777777777,
      "note": "geoparse counts all samples; panel keeps labeled subset - agreement expected to be
>="
    }
  ],
  "conclusion": "accession identity and sample rosters independently confirmed"
}
```

GCN re-implementation cross-check (PyTorch Geometric)

PyG GCN AUC 0.7240 vs diagbench GCN AUC 0.6720 on the same cleveland graph and split (absolute difference 0.0520). implementations diverge - investigated

Differential-expression validation (TCGA-BRCA)

Method: per-gene Mann-Whitney U on $\log_2(\text{count}+1)$, Benjamini-Hochberg FDR. pydeseq2 GLM fit could not complete on this hardware (documented); non-parametric screen reported instead

789 of 1000 tested genes significant at $\text{padj} < 0.05$ (200 tumor vs 113 normal). widespread significant DE confirms the tumor/normal signal the BRCA classifier (AUC 0.9993) exploits

Gene	log2 fold-change	padj
ENSG00000167434	-1.83	3.78e-44
ENSG00000091704	-1.61	1.43e-43
ENSG00000248869	-1.64	3.67e-43
ENSG00000034971	-2.05	1.11e-42
ENSG00000153822	-1.47	1.87e-41
ENSG00000253434	-1.89	2.18e-41
ENSG00000181234	-1.22	5.42e-41
ENSG00000232079	-1.46	5.97e-41
ENSG00000278928	-1.66	7.84e-41
ENSG00000165478	-1.67	1.17e-40
ENSG00000224958	-1.64	1.93e-40
ENSG00000187758	-1.68	2.06e-40
ENSG00000171819	-1.60	2.73e-40
ENSG00000250961	-1.55	7.68e-40
ENSG00000186466	-1.66	8.89e-40
ENSG00000224525	-1.79	8.89e-40
ENSG00000286322	-1.72	3.28e-39
ENSG00000246022	-1.74	5.76e-39
ENSG00000236333	-1.58	6.86e-39
ENSG00000277737	-1.94	7.04e-39

Hyperparameter search bound (optuna, wdbc MLP)

Best 4-fold CV AUC after 15 TPE trials: 0.9961 (committed logistic-regression baseline: 0.9929).
Best parameters: {"h1": 67, "h2": 36, "lr": 0.00011238623122384506, "alpha": 0.002954894558726684}.

CNN saliency (captum Integrated Gradients)

Integrated-Gradients attribution on 8 pathmnist test images (IntegratedGradients (32 steps) on CNN retrained ~ 1 epoch, 8 test images); see figure `captum_pathmnist`.

Negative results

Several hypotheses did not survive contact with the protocol, and they are recorded here with the same rigor as the positive results.

1. The GAT, despite its attention mechanism, is the weakest graph model on three of four core datasets (e.g. 0.8613 $\pm$ 0.1061 on Parkinson's) - on small cohorts the attention parameters buy variance, not signal.


2. The 1D-CNN over the feature axis never wins a core dataset; feature order carries no spatial structure for it to exploit, and it is strictly dominated by the MLP it generalizes.
3. Cross-disease bridges add nothing on WDBC (+0.0004 mean delta, sign-changing across seeds): transfer cannot lift a cohort already at its noise ceiling.
4. On wpbc (prognostic breast cancer, 198 cases) no model exceeds 0.74 AUC; the published difficulty of this dataset is real and is not an artifact of weak baselines.

Each of these is preserved with full per-seed numbers in the committed result files.


Discussion

Three conclusions follow. First, honest re-benchmarking matters: under a strict protocol, the deep architectures cluster with, not above, strong classical baselines on saturated cohorts, and the dramatic published gaps shrink to protocol differences. Second, graph structure is the one architectural ingredient that buys something classical models cannot have: label efficiency through transduction, and a small but statistically real transfer signal between diseases. Third, falsifiable discovery claims are affordable: the bridge hypothesis was stated, tested by paired ablation, and is reported with its exact effect sizes and its failures. The toolchain - loaders, models, graph builders, evaluation, paper generator - is built to be rerun end-to-end by anyone, and every number in this paper can be traced to a committed JSON artifact.


Figures




auc ann thyroid




auc arrhythmia




auc cervical




auc cleveland




auc dermatology




auc diabetes early risk




auc echocardiogram




auc fertility




auc haberman




auc hcv




auc heart hungarian




auc heart switzerland




auc heart va




auc hepatitis




auc ilpd




auc lymphography




auc mammographic




auc messidor




auc new thyroid




auc parkinsons




auc pima




auc spectf




auc statlog heart




auc wdbc




auc wisconsin original




auc wpbc




bridge delta




captum pathmnist




fingerprint umap



label eff curves



label efficiency



panel auc by arm

October 1 evidence correction and completion boundary

Locked Model C accuracy mean 0.8958333333, 95% t interval [0.8183317307, 0.9733349360], fails 0.90 hypothesis and 0.928 comparator gates. Seed37 AUC 0.9628205128 and Brier 0.0679723889 do not substitute for accuracy. Post-hoc seed37 ECE 0.0468606536 is descriptive only; seeds 11/23 ECE unavailable. Malaria R-M2 from-scratch ResNet18 mean 0.9471001905, interval [0.9311493743, 0.9630510067], fails > 0.957. No locked thresholds were changed. All prior bridge significance language is exploratory: independent fair signs were not established, and five of five positive is two-sided 0.0625, not 0.031. The 50-page length does not certify dataset/tool census, novelty, or a benchmark win. Main.tex is a separate short draft, not the sole paper artifact.