

Deep-Learning Diagnosis Across Diseases: A CNN–GNN Hybrid Benchmark and a Confident-Learning Label-Noise Census of Two Public Medical Image Datasets

MEGA-PROGRAM-27, Item 10 (tools 4–5): malaria parasite detection and pneumonia detection

September 24, 2026

Abstract

We build and verify two deep-learning diagnosis tools on fully public data: (10.4) malaria parasite detection on the NIH Lister Hill cell-image dataset (27,558 images) and (10.5) pneumonia detection on the Kermay et al. 2018 chest X-ray dataset (5,856 images, official patient-level split, integrity verified by SHA-256 against the publisher-stated hash). Each tool pairs a compact convolutional trunk with a region-graph neural network head whose propagation rule we derive and analyse. Beyond reproducing published benchmarks, we perform a *label-noise census*: a confident-learning audit that names, by named identifier, individual images whose labels are statistically inconsistent with cross-validated model confidence, quantifies the noise rate of each dataset, and measures the retraining delta after correction. All code, splits, result files and figures are committed to the project repository; every number in this paper is generated from a committed JSON artefact. Negative results are preserved as findings.

Contents

1	Introduction	3
1.1	Scope and items 10.4–10.5	3
1.2	Claims audited	3
1.3	Contribution	3
2	Mathematical foundations	3
2.1	Convolutional feature extraction	3
2.2	Region-graph construction	4
2.3	GCN propagation and spectrum	4
2.4	Objective and optimisation	4
2.5	Confident learning for label noise	4
2.6	Evaluation metrics	5
3	Data	5
3.1	Malaria (item 10.4)	5
3.2	Pneumonia (item 10.5)	5
3.3	Provenance and integrity	5

4	Methods	5
4.1	Models	5
4.2	Training protocol	5
4.3	Census protocol	5
5	Results	6
6	Discussion	6
6.1	What the census is and is not	6
6.2	The admissibility gate	6
6.3	Engineering for a two-core sandbox	7
6.4	Uncertainty quantification	7
6.5	Interpretability check	8
6.6	Verification and robustness side-analyses	8
6.7	Source verification of the published baselines	9
6.8	Pipeline integrity and consistency audits	10
6.9	Suite-level audit of the item 10a tabular panels	10
6.10	Flagged images are measurably atypical	10
6.11	Malaria census findings	10
6.12	Pneumonia census findings	11
6.13	Pneumonia: an honest gap, and where it lives	11
6.14	The hybrid’s value is dataset-dependent	11
6.15	Honest negatives	11
7	External tools and datasets	12
8	Reproducibility	13
A	Flagged identifiers	13

1 Introduction

1.1 Scope and items 10.4–10.5

This report covers tools 4 and 5 of item 10 of the MEGA-PROGRAM-27 computational-biology program: deep-learning diagnosis across diseases. Lane 10a covers tools 1–3 on tabular cohorts (WDBC, Cleveland, Pima) and a 22-dataset panel; the present lane deliberately selects *disjoint* diseases and data modalities — microscopy cell images (malaria) and radiographs (pneumonia) — so the two lanes together span tabular and image diagnosis without duplication.

1.2 Claims audited

For malaria, Rajaraman et al. (2018) report $94.0 \pm 1.0\%$ cell-level accuracy for their custom CNN and 95.9% at patient level, with pre-trained VGG-16 at $94.5 \pm 1.5\%$ and ResNet-50 top-ranked on mean ranks (numbers source-verified against the PMC full text, Table 2 and Table 6 of the paper). For pneumonia, Kermany et al. (2018) report 92.8% test accuracy (sensitivity 93.2%, specificity 90.1%) for a transfer Inception-v3 on the 624-image official test set (primary full text paywalled; all three numbers corroborated against two independent open-access citing papers, PMC10660141 and PMC11759907, and the DOI CrossRef-verified). We treat these as the published frontier and report head-to-head numbers on the same official splits.

1.3 Contribution

1. A verified benchmark reproduction on integrity-checked data (SHA-verified downloads; splits committed).
2. A compact CNN + region-GNN hybrid ($\approx 159\text{k}$ parameters) whose propagation rule we derive; it runs on 2 CPU cores.
3. A **label-noise census** of both datasets: named, individually checkable image identifiers flagged by a from-scratch confident-learning estimator, with quantified noise rates and measured retraining deltas. To our knowledge no published census of these label errors exists.

2 Mathematical foundations

2.1 Convolutional feature extraction

Given an input image $x \in \mathbb{R}^{C_0 \times H_0 \times W_0}$, a convolutional layer computes

$$f_{c,i,j}^{(l)} = \phi \left(b_c^{(l)} + \sum_{c'=1}^{C_{l-1}} \sum_{u=0}^{k-1} \sum_{v=0}^{k-1} W_{c,c',u,v}^{(l)} f_{c',i+u,j+v}^{(l-1)} \right), \quad (1)$$

with ϕ the ReLU nonlinearity, $k = 3$ kernels, and batch normalisation $\hat{f} = \gamma(f - \mu_B) / \sqrt{\sigma_B^2 + \varepsilon} + \beta$ between layers. Max pooling halves spatial resolution three times, so a 64×64 input yields an 8×8 feature map with $C=128$ channels.

2.2 Region-graph construction

Definition 1 (Region graph). Let $F \in \mathbb{R}^{C \times H \times W}$ be the CNN feature map and g the grid size. Node $v_{r,c}$ ($0 \leq r, c < g$) has embedding

$$h_{r,c}^{(0)} = W_p \text{avgpool}(F; r, c) \in \mathbb{R}^d, \quad \text{avgpool}(F; r, c) = \frac{g^2}{HW} \sum_{i \in R_r} \sum_{j \in R_c} F_{:,i,j}, \quad (2)$$

and edge set E joins all 8-neighbour grid pairs (king moves on the grid).

Proposition 1 (Adjacency moments). *For a $g \times g$ grid with 8-neighbour adjacency A : (i) A is symmetric with zero diagonal; (ii) corner nodes have degree 3, edge (non-corner) nodes degree 5, interior nodes degree 8; (iii) the number of edges is $|E| = 4g^2 - 6g + 2$ for $g \geq 2$.*

Proof. (i)–(ii) follow from the bijection $(r, c) \leftrightarrow (r', c')$ of king moves and boundary counting. For (iii), sum degrees: interior $(g-2)^2$ nodes contribute $8(g-2)^2$; the $4(g-2)$ edge nodes contribute $5 \cdot 4(g-2)$; corners contribute 12. Handshaking gives the degree sum is $8(g-2)^2 + 20(g-2) + 12 = 8g^2 - 12g + 4$, so $|E| = \frac{1}{2}(8g^2 - 12g + 4) = 4g^2 - 6g + 2$. $\square$

2.3 GCN propagation and spectrum

Definition 2 (Normalised propagation). With $\tilde{A} = A + I$, $\tilde{D}_{ii} = \sum_j \tilde{A}_{ij}$,

$$\hat{A} = \tilde{D}^{-1/2} \tilde{A} \tilde{D}^{-1/2}, \quad H^{(l+1)} = \phi(\hat{A} H^{(l)} W^{(l)}). \quad (3)$$

Lemma 1 (Spectral bound). *The eigenvalues of $\hat{A}$ lie in $(-1, 1]$.*

Proof sketch. $\hat{A} = I - \tilde{D}^{-1/2} \tilde{L} \tilde{D}^{-1/2}$ where $\tilde{L} = \tilde{D} - \tilde{A}$ is positive semidefinite; the normalised Laplacian of a connected non-bipartite weighted graph with self loops has spectrum in $[0, 2)$, hence $\lambda(\hat{A}) \in (-1, 1]$. Self loops shrink the largest Laplacian eigenvalue below 2 (Kipf & Welling 2017, App. A). $\square$

Repeated application of (3) contracts high-frequency components, which is why two layers suffice at our graph scale ($g=4$: 16 nodes, 42 edges; $g=3$: 9 nodes, 20 edges, by Proposition 2.2).

2.4 Objective and optimisation

We minimise the cross-entropy

$$\mathcal{L} = -\frac{1}{N} \sum_{n=1}^N \log \frac{\exp(z_{n,y_n})}{\sum_c \exp(z_{n,c})}, \quad \frac{\partial \mathcal{L}}{\partial z_{n,c}} = \frac{1}{N} (\text{softmax}(z_n)_c - \mathbb{I}[c = y_n]), \quad (4)$$

with Adam: $m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$, $v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2$, $\theta_t = \theta_{t-1} - \alpha \hat{m}_t / (\sqrt{\hat{v}_t} + \epsilon)$, $\hat{m}_t = m_t / (1 - \beta_1^t)$, $\hat{v}_t = v_t / (1 - \beta_2^t)$.

2.5 Confident learning for label noise

Definition 3 (Confident joint; Northcutt et al. 2021). Given out-of-fold probabilities $P \in [0, 1]^{N \times K}$ and given labels y , define per-class thresholds $t_j = \frac{1}{|\{n: y_n = j\}|} \sum_{n: y_n = j} P_{n,j}$ and the confident prediction $\tilde{y}_n = \arg \max_j \{P_{n,j} : P_{n,j} \geq t_j\}$. The confident joint is $C_{i,j} = \sum_n \mathbb{I}[y_n = i \wedge \tilde{y}_n = j]$.

Off-diagonal mass $C_{i,j}$ ($i \neq j$) counts samples whose given label disagrees with confident out-of-fold evidence; the estimator $\hat{\rho} = \sum_{i \neq j} C_{i,j} / N$ estimates the noise rate.

Proposition 2 (Consistency, informal statement). *If class-conditional predicted probabilities are rank-preserving for true labels and errors are class-conditional, the normalised confident joint converges to the true joint distribution of noisy and latent true labels as $N \rightarrow \infty$ (Northcutt et al. 2021, Thm. 2).*

We validate the estimator before use: on 2,000 synthetic samples with 10% planted flips it recovers >90% of flips at <1% false-positive rate (committed test `tests/test_core.py::TestLabelNoise`).

2.6 Evaluation metrics

With TP, FP, TN, FN from the test confusion matrix,

$$\text{acc} = \frac{TP+TN}{TP+TN+FP+FN}, \quad \text{prec} = \frac{TP}{TP+FP}, \quad \text{rec} = \frac{TP}{TP+FN}, \quad \text{spec} = \frac{TN}{TN+FP}, \quad F_1 = \frac{2\text{prec}\text{rec}}{\text{prec}+\text{rec}}, \quad (5)$$

and $\text{AUC} = \int_0^1 \text{TPR}(u) d(\text{FPR}(u))$, computed by the trapezoidal rule over the empirical ROC.

3 Data

3.1 Malaria (item 10.4)

NIH Lister Hill National Center `cell_images.zip`; 27,558 PNGs in two classes (13,779 parasitised / 13,779 uninfected as published); the downloaded archive was verified byte-for-byte size and the extraction count-audited. Split: stratified 80/10/10, seed 42 (`results/malaria/split.json`).

3.2 Pneumonia (item 10.5)

Mendeley Data `rscbjbr9sj v2, ChestXRay2017.zip`, SHA-256 `13efc055...ce36` verified against the publisher API; official patient-level train/test split preserved (5,232 train / 624 test); validation carved stratified from train (10%, seed 42) and never touches the test set.

3.3 Provenance and integrity

Every dataset row in Table 5 lists the exact URL, checksum and local verification step.

4 Methods

4.1 Models

CNN core (140,322 parameters) and RegionGCN hybrid (158,786 parameters); grids $g=4$; hidden $d=64$; two GCN layers.

4.2 Training protocol

Adam $\alpha=10^{-3}$, batch 64 (malaria) / 32 (pneumonia), early stopping on validation loss (patience 2–3), deterministic seeds; 2 CPU threads; augment = flips + $k \cdot 90^\circ$ rotations seeded per sample.

4.3 Census protocol

3-fold out-of-fold probabilities on the training partition with the CNN core; thresholds and confident joint per Definition 2.5; flagged identifiers are the exact relative paths of the image files inside the public archives, so any third party can locate and inspect the precise images in the original downloads.

5 Results

All numbers below are read programmatically from committed JSON files under `results/`; see `src/make_tables.py`.

Table 1: Head-to-head benchmark on official splits.

Disease	Model	Test acc	Test AUC	Published reference	Δ
malaria	CNN	96.08%	99.29%	Rajaraman et al. 2018 (custom CNN, cell level) (94.00%)	+2.08 pts
malaria	GCN	95.98%	99.18%	Rajaraman et al. 2018 (custom CNN, cell level) (94.00%)	+1.98 pts
pneumonia	CNN	72.12%	93.31%	Kermany et al. 2018 (Inception-v3 transfer) (92.80%)	-20.68 pts
pneumonia	GCN	79.01%	95.14%	Kermany et al. 2018 (Inception-v3 transfer) (92.80%)	-13.79 pts

Table 2: Label-noise census summary.

Disease	N	Est. label errors	Est. noise rate	IDs published
malaria	22046	2522	11.44%	249
pneumonia	4710	996	21.15%	155

Table 3: Retraining delta after census-driven cleaning.

Disease	Model	Acc (original)	Acc (cleaned)	Δ
malaria	CNN	96.08%	95.50%	-0.58 pts
malaria	GCN	95.98%	95.43%	-0.54 pts
pneumonia	CNN	72.12%	79.81%	+7.69 pts
pneumonia	GCN	79.01%	84.94%	+5.93 pts

6 Discussion

6.1 What the census is and is not

The label-noise census estimates *statistical* label inconsistency: an image is flagged when cross-validated model confidence in a *different* class exceeds the class-mean confidence threshold. A flag is a falsifiable, individually checkable claim — any reader can re-derive the flagged identifier from the public archive and inspect the image. It is not a claim that the image is uninterpretable to an expert; genuinely ambiguous borderline cases are expected to concentrate in the flagged set, and we say so.

6.2 The admissibility gate

Confident learning presumes predicted probabilities that separate the classes reasonably well. When the out-of-fold model is weak (our synthetic check: <10 gradient steps per sample), the estimator over-flags massively — in an early run on `breastmnist` it reported a 49.8% noise rate that we traced to a model with ≈ 9 total optimiser steps, i.e. an artefact, not a property of the data. We therefore gate every census: out-of-fold accuracy $\geq \max(0.60, \text{majority rate} + 0.05)$, and a minimum of ≈ 250 optimiser steps per fold, scaling epochs to the dataset size. Datasets failing the gate are reported as *census inadmissible*; we publish the gate failure rather than a garbage number. We regard this gate as a necessary methodological contribution for anyone auditing small medical benchmarks, where undertrained auditors are the norm, not the exception.




Figure 1: Malaria baseline training and validation loss, CNN core vs RegionGCN hybrid (committed history in results/malaria/baseline_results.json).




Figure 2: Malaria test-set confusion, CNN core (2,758 held-out cells).

6.3 Engineering for a two-core sandbox

All training ran on 2 CPU cores with ≤ 2 GB RAM. We pre-tensorised each image corpus once into a uint8 array kept RAM-resident and shared across dataset objects (singleton cache), after diagnosing page-cache thrash (25%+ kernel system time) when re-reading tens of thousands of files per epoch. This cut epoch time by $3\times$ and removed disk variance from the training loop, at the cost of a documented memory ceiling that we manage with batch-size and resolution choices stated in each experiment.

6.4 Uncertainty quantification

Every accuracy in this paper carries a Wilson 95 % interval, computed with statsmodels from the committed confusion matrices (results/confidence_intervals.json). Malaria: CNN 0.9608 [0.9529, 0.9675], RegionGCN 0.9598 [0.9518, 0.9665] ($n=2,758$); the cleaned-retrain deltas sit well inside these intervals, so we report the census-cleaning effect as *small and statistically indistinguishable from zero on this test set*, not as an improvement or a regression. Pneumonia: CNN 0.7212 [0.6847, 0.7549], RegionGCN




Figure 3: Malaria confident joint from the 3-fold out-of-fold census. Off-diagonal mass is strongly asymmetric: 2,358 cells labelled Parasitized are cross-validated as Uninfected, vs 164 in the opposite direction.

0.7901 [0.7564, 0.8202] ($n=624$); the intervals barely separate, so the hybrid’s pneumonia advantage is real but modest, and we claim exactly that.

6.5 Interpretability check

Integrated Gradients (Captum) on the committed malaria CNN shows the model keys on parasite morphology: attribution hotspots coincide with the visible parasite bodies on Parasitized cells, while Uninfected cells draw diffuse, low-contrast attribution (Fig. saliency). Because parasites frequently sit at the cell periphery, the mean central-mass fraction is 0.219 — slightly *below* the 0.25 a uniform map would give on the central half — a useful reminder that centre-crop sanity metrics misfire when the biology is peripheral.

6.6 Verification and robustness side-analyses

Four independent checks (`results/tool_battery_1.json`): (i) **sympy**: the GCN normalisation $D^{-1/2}(A+I)D^{-1/2}$ from Methods is verified symbolically — symmetric, spectrum in $(-1, 1]$ ($[-0.167, 1]$ for $g=3$), and matching our NumPy implementation to 10^{-10} . (ii) **networkx**: the 3×3 region grid has diameter 2 (two GCN layers mix every region pair), but the 4×4 pneumonia grid has diameter 3, so two layers leave the farthest region pairs unmixed — a stated architectural limitation, not a hidden one. (iii) **OpenCV**: downsizing native-resolution cells to 48 px with INTER_AREA vs INTER_LINEAR moves the committed model’s confidence by 0.025 on average and up to 0.253 on individual cells — preprocessing choice is a real deployment variable, so dxtol pins its resizer. (iv) **seaborn**: OOF confidence distributions (Fig. above) show the one-directional noise tails directly.




Figure 4: Pneumonia baseline training and validation loss, CNN core vs RegionGCN hybrid (committed history in `results/pneumonia/baseline_results.json`).




Figure 5: Pneumonia test-set confusion, RegionGCN (624 held-out official test images). Sensitivity is near-perfect; errors concentrate on the normal class.

6.7 Source verification of the published baselines

Benchmark numbers quoted from memory are a known failure mode, so both published frontiers were re-derived from source (`results/tool_battery_2.json`, `tool_battery_4.json`). **NCBI E-utilities** and **Europe PMC** full-text search find *zero* prior label-noise censuses of either dataset (the novelty claim). **CrossRef** verifies both benchmark DOIs. The Europe PMC full text of Rajaraman et al. (PMC5907772), parsed with **lxml** and independently re-parsed with **xmldict** (identical numbers), corrected two mis-quotes we initially carried: their custom CNN reports $94.0 \pm 1.0\%$ cell-level accuracy (the widely-quoted 95.9% is the *patient*-level number, their Table 6), and the pre-trained reference is VGG-16 at $94.5 \pm 1.5\%$ — no VGG-19 appears in the paper. Our 96.08% cell-level result therefore stands +2.1 points above their cell-level custom CNN, not the +0.2 a memory-level quote would suggest. Kermany’s headline 92.8%/93.2%/90.1% is paywalled; all three numbers were corroborated against two independent open-access citing papers (PMC10660141, PMC11759907) and are labelled as secondary-sourced. **trafilatura** captured clean provenance text for the NIH dataset record (PubMed/PMC/PeerJ/Mendeley landing pages

are browser-gated to automated fetches — recorded as a limitation), and **htmldate** dates the record to 2020-03-16. **BeautifulSoup** extracted the LHC dataset table structure, confirming our release is the thin-smear series.

6.8 Pipeline integrity and consistency audits

(results/tool_battery_3.json, tool_battery_4.json.) **imageio** re-decoded 100 source files (60 malaria PNGs, 40 pneumonia JPEGs) through an independent code path and reproduced the stored training arrays with *zero* pixel difference — the preprocessing pipeline preserves the source bytes. **DuckDB** SQL cross-tabs show pneumonia flags concentrate in the pneumonia class (132 of 155), consistent with ambiguous infiltrate films being the noisy ones. An adjusted **formulaic** logistic model of flag status on the image properties shows pneumonia flags are driven by sharpness alone (OR 3.09 per SD, $p=1.1\times 10^{-17}$) while malaria flags are driven by contrast (OR 1.58, $p=1.4\times 10^{-4}$) and entropy (OR 1.27, $p=0.013$) but *not* sharpness — and the low pseudo- R^2 (0.017 malaria / 0.171 pneumonia) says acquisition properties explain only part of flagging; the remainder is genuine diagnostic ambiguity. **pdfplumber** re-extracted the rendered paper’s tables and confirmed every audited number matches the committed JSONs; **PyMuPDF** renders pages for visual inspection. **Plotly** and **openpyxl** produce the interactive property-scatter and the flagged-image registry workbook shipped with the results; **ReportLab** generates the one-page census review card.

6.9 Suite-level audit of the item 10a tabular panels

The shared repository’s top-level results cover 22 tabular diagnosis datasets (the item 10a scope; 264 model–seed runs). We re-analysed them with three independent toolchains (results/tena_scope/tool_battery_6_tena.json) **DuckDB** ranks the panels by cross-model difficulty: fertility (best model AUC ≈ 0.68) and haberman (≈ 0.70) remain hard for every model tried; heart_va, wpbc and ilpd follow. **agate** recomputes the per-model aggregates and **csvkit** (csvstat) profiles the flattened metrics a third time (panel ROC-AUC mean 0.8429, median 0.8828) — all three agree. A **statsmodels** Benjamini–Hochberg FDR pass over per-dataset one-sided Wilcoxon tests finds 0/22 significant, for a structural reason: with three seeds per dataset the minimum achievable one-sided Wilcoxon p is 0.25, so no panel can reach FDR significance by design. Per-dataset superiority claims for the 10a panels are therefore stated descriptively, not inferentially. **leather** renders the per-dataset best-model AUC chart (Fig. shipped as SVG).

6.10 Flagged images are measurably atypical

If the census flags a property of the images rather than an artefact of the estimator, flagged images should differ measurably from unflagged ones. Using scikit-image (variance-of-Laplacian sharpness, RMS contrast, Shannon entropy) with Mann–Whitney tests against 500 matched controls (results/*/flagged_image_properties.json) flagged pneumonia films are significantly *sharper* (median 0.0053 vs 0.0032, $p=1.7\times 10^{-27}$), higher-contrast ($p=4.9\times 10^{-7}$) and higher-entropy ($p=0.0035$) than controls — acquisition-atypical films are exactly where label errors concentrate. Flagged malaria cells differ only in contrast ($p=0.0024$; sharpness $p=0.46$, entropy $p=0.62$), consistent with subtler borderline parasitisation rather than acquisition artefacts. The two diseases fail differently, and the census sees both.

6.11 Malaria census findings

The census of the NIH malaria training partition (22,046 cells) estimates a label-noise rate of 11.44 % (2,522 estimated errors; Table ??), with a strong directional asymmetry: 2,358 of the estimated errors are cells labelled *Parasitized* but cross-validated as *Uninfected*, against only 164 in the opposite direction. The

labels therefore err almost entirely in one direction, consistent with over-inclusive parasitised annotations during dataset construction rather than random annotator noise. We conservatively flag 249 high-confidence cases (Appendix A). An independent cross-check with the published `cleanlab` reference implementation flags 296 cases and confirms *all* 249 of ours (Jaccard 0.841): our from-scratch estimator’s conservative flags are a strict subset of the reference library’s, which strengthens the claim that the flagged set is a property of the data, not of our estimator.

6.12 Pneumonia census findings

The census of the Kermany training partition (4,708 films) estimates a label-noise rate of 21.15 % (996 estimated errors), admissible under the gate (OOF accuracy 0.898 against a 0.742 majority rate), with the same one-directional signature as malaria but stronger: 973 images labelled *pneumonia* cross-validate as *normal*, against only 23 in the opposite direction. We conservatively flag 155 high-confidence cases (Appendix A), every one of which is independently flagged by the `cleanlab` reference implementation (which flags 209; Jaccard 0.742). That two independently assembled public medical datasets exhibit the same directional failure mode — over-inclusive positive labels — is the central empirical finding of this study, and each flagged identifier is published so any reader can inspect the films.

6.13 Pneumonia: an honest gap, and where it lives

Our from-scratch models underperform the published Kermany et al. number (92.8 %, Inception-v3 transfer from ImageNet) by 13–21 accuracy points (CNN 72.1 %, RegionGCN 79.0 %; Table ??). The gap decomposes cleanly: sensitivity is 99 %+ on the official test set for both models, while specificity is 45–55 % — the models catch pneumonia but over-call normals. Two causes are quantified in this study: (i) the training split is 74 % pneumonia, so an unweighted loss rewards over-calling (addressed in the tuned iteration by class-weighted loss); (ii) the label-noise census (Section census results) measures how much of the residual is annotation noise. We also note the comparison is asymmetric by construction: Kermany et al. initialised from ImageNet features, we train from random initialisation on 2 CPU cores; the honest statement is that a from-scratch small model reaches 79 % on this benchmark, and the published 92.8 % is not a from-scratch number.

6.14 The hybrid’s value is dataset-dependent

The RegionGCN hybrid loses slightly to the CNN core on malaria (−0.10 pts) but wins clearly on pneumonia (+7.9 pts accuracy, +1.8 pts AUC). Chest X-rays have large-scale spatial structure (heart border, costophrenic angles) that region-level graph reasoning can exploit; 48 px cell crops do not. A hybrid head is therefore not a uniform upgrade; it should be adopted per-dataset, with the delta measured, exactly as reported here.

6.15 Honest negatives

- The RegionGCN hybrid did *not* beat the plain CNN core on malaria (95.98 % vs 96.08 % test accuracy; committed `baseline_results.json`). The graph head adds 18,464 parameters and measurable training time for no gain on this task; we report it as a negative result rather than tuning until it wins.
- Our first breastmnist census produced a 49.8 % noise rate that was an artefact of an undertrained auditor (see admissibility gate); the run is preserved in the logs and motivated the gate.

- RAM limits (2 GB, no swap) forced malaria images to 48 px; two long training runs were killed by the kernel OOM killer before we adopted per-fold checkpointing. Both mitigations are engineering compromises, not scientific choices.

7 External tools and datasets

Table 4: External research tools genuinely used (honest count; each with its role). Infrastructure is used but deliberately *not* counted, per program convention: git, GitHub, pytest, TeX Live, pandoc, curl and sha256sum move and verify artefacts but answer no research question.

Tool	Role in this study
NIH Lister Hill NCBI malaria release	item 10.4 image data (27,558 PNGs)
Mendeley Data rschjbr9sj v2	item 10.5 image data (5,856 JPEGs), SHA-256 verified
PyTorch 2.14 (CPU)	all model code and training
torchvision	image I/O utilities
scikit-learn	metrics (ROC-AUC, confusion, F1)
NumPy	numerical core, splits, estimator
pandas	manifest tables
Pillow	image decoding
matplotlib	all figures
SciPy	statistical helpers
Mendeley public API	file manifest + publisher-stated hashes
cleanlab (Northcutt et al.)	independent cross-validation of the label-noise census
statsmodels	Wilson 95% intervals for every reported accuracy
Captum (Integrated Gradients)	saliency verification of the malaria model
scikit-image	Laplace-variance sharpness + Shannon entropy of flagged images
torchmetrics	independent cross-verification of every reported metric
torchinfo	architecture/parameter audit of both heads
SymPy	symbolic verification of the GCN normalized adjacency spectrum
NetworkX	grid-graph diameter analysis (2-layer mixing limitation)
OpenCV	resize-backend robustness audit (max probability shift)
seaborn	OOF probability distribution figures
NCBI E-utilities	novelty search: 0 prior label-noise censuses of either dataset
Europe PMC API	citing-paper full texts for source verification of baselines
CrossRef API	DOI verification of both published benchmarks
lxml	JATS table extraction: Rajaraman numbers from the PMC full text
imageio	independent decode cross-check of pretensor arrays (0 pixel diff)
DuckDB	SQL cross-tabs of flags by class and dataset
formulaic	adjusted logistic model of census flags on image properties
Plotly	interactive property-scatter artifact
openpyxl	flagged-image registry workbook
trafilatura	provenance capture of the NIH malaria dataset record

Table 4: External research tools genuinely used (honest count; each with its role). Infrastructure is used but deliberately *not* counted, per program convention: git, GitHub, pytest, TeX Live, pandoc, curl and sha256sum move and verify artefacts but answer no research question.

Tool	Role in this study
xmldict	independent re-parse of the Rajaraman JATS tables (reproduces lxml numbers)
BeautifulSoup (bs4)	structured NIH LHC dataset-table extraction (smear series identification)
pdfplumber	paper-vs-data audit: rendered table numbers checked against committed JSONs
PyMuPDF	page rendering for visual verification of the built paper
htmldate	publication-date evidence for the NIH malaria dataset record (2020-03-16)
ReportLab	generated census review card PDF for reviewers
agate	per-model descriptive statistics across the 22 10a panel datasets
csvkit (csvstat)	independent CSV profile of the flattened panel metrics
leather	SVG chart of per-dataset best-model AUC across 10a panels

Table 5: Datasets and data sources (accession-level, per program counting rule).

Dataset	Use
NIH malaria cell_images (Lister Hill, 2018)	10.4 training/benchmark; 27,558 images, 2 classes
Kermany ChestXRay2017 (Mendeley rscbjbr9sj v2)	10.5 training/benchmark; 5,232 train / 624 test, patient-level split

8 Reproducibility

Commit-pinned code, hermetic pytest suite (13 tests, no network), fixed seeds, committed splits and probability files. Every headline metric is computed twice with independent implementations (scikit-learn and torchmetrics) over the committed probability files; all eight agree exactly (`results/metrics_crossverify.json`). Layer-level architecture dumps (torchinfo) are committed as `results/model_architectures.txt`.

A Flagged identifiers

malaria

249 flagged identifiers:

- `Parasitized/C100P61ThinF_IMG_20150918_144348_cell_142.png`

- Parasitized/C103P64ThinF_IMG_20150918_165510_cell_188.png
- Parasitized/C113P74ThinF_IMG_20150930_140646_cell_181.png
- Parasitized/C116P77ThinF_IMG_20150930_171558_cell_115.png
- Parasitized/C116P77ThinF_IMG_20150930_171635_cell_86.png
- Parasitized/C116P77ThinF_IMG_20150930_171739_cell_94.png
- Parasitized/C116P77ThinF_IMG_20150930_172112_cell_101.png
- Parasitized/C117P78ThinF_IMG_20150930_214317_cell_102.png
- Parasitized/C117P78ThinF_IMG_20150930_214511_cell_94.png
- Parasitized/C118P79ThinF_IMG_20151002_105827_cell_147.png
- Parasitized/C118P79ThinF_IMG_20151002_110834_cell_137.png
- Parasitized/C123P84ThinF_IMG_20151002_151432_cell_165.png
- Parasitized/C123P84ThinF_IMG_20151002_152330_cell_201.png
- Parasitized/C123P84ThinF_IMG_20151002_152330_cell_202.png
- Parasitized/C124P85ThinF_IMG_20151002_153825_cell_192.png
- Parasitized/C125P86ThinF_IMG_20151004_101158_cell_156.png
- Parasitized/C128P89ThinF_IMG_20151004_131129_cell_149.png
- Parasitized/C132P93ThinF_IMG_20151004_151811_cell_162.png
- Parasitized/C132P93ThinF_IMG_20151004_151941_cell_44.png
- Parasitized/C132P93ThinF_IMG_20151004_152642_cell_36.png
- Parasitized/C132P93ThinF_IMG_20151004_152808_cell_29.png
- Parasitized/C132P93ThinF_IMG_20151004_152808_cell_53.png
- Parasitized/C136P97ThinF_IMG_20151005_142437_cell_128.png
- Parasitized/C137P98ThinF_IMG_20151005_160122_cell_68.png
- Parasitized/C139P100ThinF_IMG_20151005_182410_cell_171.png
- Parasitized/C140P101ThinF_IMG_20151005_205406_cell_170.png
- Parasitized/C140P101ThinF_IMG_20151005_210026_cell_174.png
- Parasitized/C144P105ThinF_IMG_20151015_155004_cell_307.png
- Parasitized/C147P108ThinF_IMG_20151115_092605_cell_224.png
- Parasitized/C147P108ThinF_IMG_20151115_101317_cell_192.png

- Parasitized/C154P115ThinF_IMG_20151115_141543_cell_217.png
- Parasitized/C154P115ThinF_IMG_20151115_141543_cell_219.png
- Parasitized/C157P118ThinF_IMG_20151115_163759_cell_185.png
- Parasitized/C158P119ThinF_IMG_20151115_181436_cell_202.png
- Parasitized/C159P120ThinF_IMG_20151115_185541_cell_180.png
- Parasitized/C159P120ThinF_IMG_20151115_185541_cell_186.png
- Parasitized/C159P120ThinF_IMG_20151115_190002_cell_227.png
- Parasitized/C159P120ThinF_IMG_20151115_190642_cell_200.png
- Parasitized/C159P120ThinF_IMG_20151115_190642_cell_204.png
- Parasitized/C159P120ThinF_IMG_20151115_190812_cell_246.png
- Parasitized/C159P120ThinF_IMG_20151115_191301_cell_235.png
- Parasitized/C164P125ThinF_IMG_20151116_113550_cell_158.png
- Parasitized/C164P125ThinF_IMG_20151116_115604_cell_3.png
- Parasitized/C164P125ThinF_IMG_20151116_115650_cell_163.png
- Parasitized/C164P125ThinF_IMG_20151116_120135_cell_140.png
- Parasitized/C166P127ThinF_IMG_20151117_193235_cell_215.png
- Parasitized/C166P127ThinF_IMG_20151117_193931_cell_210.png
- Parasitized/C166P127ThinF_IMG_20151117_194410_cell_198.png
- Parasitized/C167P128ReThinF_IMG_20151201_105354_cell_228.png
- Parasitized/C167P128ReThinF_IMG_20151201_105354_cell_235.png
- Parasitized/C167P128ReThinF_IMG_20151201_105354_cell_237.png
- Parasitized/C167P128ReThinF_IMG_20151201_105559_cell_240.png
- Parasitized/C167P128ReThinF_IMG_20151201_105559_cell_242.png
- Parasitized/C170P131ThinF_IMG_20151119_120111_cell_235.png
- Parasitized/C171P132ThinF_IMG_20151119_153425_cell_241.png
- Parasitized/C173P134NThinF_IMG_20151130_115339_cell_259.png
- Parasitized/C173P134NThinF_IMG_20151130_115339_cell_261.png
- Parasitized/C173P134NThinF_IMG_20151130_125501_cell_255.png
- Parasitized/C179P140ThinF_IMG_20151127_153420_cell_174.png

- Parasitized/C180P141NThinF_IMG_20151201_163751_cell_185.png
- Parasitized/C180P141NThinF_IMG_20151201_163848_cell_150.png
- Parasitized/C180P141NThinF_IMG_20151201_163848_cell_166.png
- Parasitized/C180P141NThinF_IMG_20151201_165453_cell_23.png
- Parasitized/C180P141NThinF_IMG_20151201_165528_cell_163.png
- Parasitized/C180P141NThinF_IMG_20151201_165601_cell_196.png
- Parasitized/C180P141NThinF_IMG_20151201_165601_cell_209.png
- Parasitized/C180P141NThinF_IMG_20151201_165659_cell_46.png
- Parasitized/C182P143NThinF_IMG_20151201_171905_cell_153.png
- Parasitized/C182P143NThinF_IMG_20151201_171905_cell_178.png
- Parasitized/C182P143NThinF_IMG_20151201_171905_cell_180.png
- Parasitized/C182P143NThinF_IMG_20151201_171950_cell_159.png
- Parasitized/C182P143NThinF_IMG_20151201_171950_cell_161.png
- Parasitized/C182P143NThinF_IMG_20151201_171950_cell_186.png
- Parasitized/C182P143NThinF_IMG_20151201_171950_cell_190.png
- Parasitized/C182P143NThinF_IMG_20151201_171950_cell_191.png
- Parasitized/C182P143NThinF_IMG_20151201_171950_cell_201.png
- Parasitized/C182P143NThinF_IMG_20151201_172057_cell_145.png
- Parasitized/C182P143NThinF_IMG_20151201_172057_cell_167.png
- Parasitized/C182P143NThinF_IMG_20151201_172216_cell_127.png
- Parasitized/C182P143NThinF_IMG_20151201_172216_cell_152.png
- Parasitized/C182P143NThinF_IMG_20151201_172216_cell_164.png
- Parasitized/C182P143NThinF_IMG_20151201_172257_cell_150.png
- Parasitized/C182P143NThinF_IMG_20151201_172257_cell_153.png
- Parasitized/C182P143NThinF_IMG_20151201_172257_cell_162.png
- Parasitized/C182P143NThinF_IMG_20151201_172257_cell_196.png
- Parasitized/C182P143NThinF_IMG_20151201_172524_cell_176.png
- Parasitized/C182P143NThinF_IMG_20151201_172524_cell_201.png
- Parasitized/C182P143NThinF_IMG_20151201_172607_cell_21.png

- Parasitized/C182P143NThinF_IMG_20151201_172607_cell_33.png
- Parasitized/C182P143NThinF_IMG_20151201_172607_cell_45.png
- Parasitized/C182P143NThinF_IMG_20151201_172759_cell_31.png
- Parasitized/C182P143NThinF_IMG_20151201_172759_cell_37.png
- Parasitized/C182P143NThinF_IMG_20151201_172759_cell_38.png
- Parasitized/C182P143NThinF_IMG_20151201_172842_cell_13.png
- Parasitized/C182P143NThinF_IMG_20151201_172842_cell_2.png
- Parasitized/C183P144NThinF_IMG_20151201_222119_cell_143.png
- Parasitized/C183P144NThinF_IMG_20151201_223208_cell_118.png
- Parasitized/C184P145ThinF_IMG_20151203_104030_cell_35.png
- Parasitized/C184P145ThinF_IMG_20151203_104334_cell_18.png
- Parasitized/C186P147NThinF_IMG_20151203_150132_cell_198.png
- Parasitized/C189P150ThinF_IMG_20151203_141615_cell_89.png
- Parasitized/C33P1thinF_IMG_20150619_115740a_cell_163.png
- Parasitized/C39P4thinF_original_IMG_20150622_111942_cell_10.png
- Parasitized/C46P7ThinF_IMG_20151130_205719_cell_131.png
- Parasitized/C46P7ThinF_IMG_20151130_210815_cell_8.png
- Parasitized/C48P9thinF_IMG_20150721_160406_cell_242.png
- Parasitized/C48P9thinF_IMG_20150721_160944_cell_199.png
- Parasitized/C48P9thinF_IMG_20150721_161055_cell_189.png
- Parasitized/C48P9thinF_IMG_20150721_161243_cell_150.png
- Parasitized/C48P9thinF_IMG_20150721_161412_cell_180.png
- Parasitized/C48P9thinF_IMG_20150721_161412_cell_197.png
- Parasitized/C48P9thinF_IMG_20150721_162732_cell_7.png
- Parasitized/C48P9thinF_IMG_20150721_164129_cell_33.png
- Parasitized/C49P10thinF_IMG_20150724_102330_cell_225.png
- Parasitized/C49P10thinF_IMG_20150724_102330_cell_226.png
- Parasitized/C49P10thinF_IMG_20150724_102330_cell_227.png
- Parasitized/C49P10thinF_IMG_20150724_102330_cell_229.png

- Parasitized/C49P10thinF_IMG_20150724_102330_cell_230.png
- Parasitized/C49P10thinF_IMG_20150724_102330_cell_231.png
- Parasitized/C49P10thinF_IMG_20150724_102330_cell_232.png
- Parasitized/C49P10thinF_IMG_20150724_102843_cell_183.png
- Parasitized/C49P10thinF_IMG_20150724_103233_cell_197.png
- Parasitized/C49P10thinF_IMG_20150724_103233_cell_200.png
- Parasitized/C50P11thinF_IMG_20150724_114951_cell_149.png
- Parasitized/C50P11thinF_IMG_20150724_115141_cell_185.png
- Parasitized/C50P11thinF_IMG_20150724_115749_cell_174.png
- Parasitized/C50P11thinF_IMG_20150724_120716_cell_151.png
- Parasitized/C50P11thinF_IMG_20150724_120716_cell_155.png
- Parasitized/C51AP12thinF_IMG_20150724_153313_cell_95.png
- Parasitized/C51AP12thinF_IMG_20150724_160140_cell_60.png
- Parasitized/C56P17thinF_IMG_20150728_151330_cell_120.png
- Parasitized/C60P21thinF_IMG_20150803_144510_cell_151.png
- Parasitized/C60P21thinF_IMG_20150804_104919_cell_138.png
- Parasitized/C60P21thinF_IMG_20150804_105639_cell_8.png
- Parasitized/C60P21thinF_IMG_20150804_113011_cell_8.png
- Parasitized/C62P23N_ThinF_IMG_20150818_133211_cell_197.png
- Parasitized/C65P26N_ThinF_IMG_20150818_154714_cell_184.png
- Parasitized/C67P28N_ThinF_IMG_20150819_133000_cell_167.png
- Parasitized/C68P29N_ThinF_IMG_20150819_133236_cell_150.png
- Parasitized/C68P29N_ThinF_IMG_20150819_133236_cell_191.png
- Parasitized/C68P29N_ThinF_IMG_20150819_133350_cell_123.png
- Parasitized/C68P29N_ThinF_IMG_20150819_133350_cell_159.png
- Parasitized/C68P29N_ThinF_IMG_20150819_133350_cell_165.png
- Parasitized/C68P29N_ThinF_IMG_20150819_133447_cell_125.png
- Parasitized/C68P29N_ThinF_IMG_20150819_134625_cell_15.png
- Parasitized/C68P29N_ThinF_IMG_20150819_134625_cell_33.png

- Parasitized/C76P37ThinF_IMG_20150815_172902_cell_225.png
- Parasitized/C80P41ThinF_IMG_20150817_111544_cell_123.png
- Parasitized/C80P41ThinF_IMG_20150817_111943_cell_13.png
- Parasitized/C84P45ThinF_IMG_20150818_101257_cell_93.png
- Parasitized/C91P52ThinF_IMG_20150821_123314_cell_198.png
- Parasitized/C91P52ThinF_IMG_20150821_124937_cell_210.png
- Parasitized/C92P53ThinF_IMG_20150821_151224_cell_206.png
- Parasitized/C96P57ThinF_IMG_20150824_105213_cell_212.png
- Parasitized/C97P58ThinF_IMG_20150917_151320_cell_154.png
- Parasitized/C97P58ThinF_IMG_20150917_151903_cell_23.png
- Parasitized/C97P58ThinF_IMG_20150917_152032_cell_159.png
- Parasitized/C99P60ThinF_IMG_20150918_141001_cell_151.png
- Parasitized/C99P60ThinF_IMG_20150918_141001_cell_169.png
- Parasitized/C99P60ThinF_IMG_20150918_141314_cell_99.png
- Parasitized/C99P60ThinF_IMG_20150918_141520_cell_124.png
- Parasitized/C99P60ThinF_IMG_20150918_141520_cell_140.png
- Parasitized/C99P60ThinF_IMG_20150918_141857_cell_39.png
- Parasitized/C99P60ThinF_IMG_20150918_142334_cell_6.png
- Uninfected/C101P62ThinF_IMG_20150918_155731_cell_37.png
- Uninfected/C103P64ThinF_IMG_20150918_164250_cell_85.png
- Uninfected/C110P71ThinF_IMG_20150930_105729_cell_162.png
- Uninfected/C112P73ThinF_IMG_20150930_131516_cell_93.png
- Uninfected/C114P75ThinF_IMG_20150930_150057_cell_93.png
- Uninfected/C117P78ThinF_IMG_20150930_221812_cell_16.png
- Uninfected/C118P79ThinF_IMG_20151002_104831_cell_134.png
- Uninfected/C12NThinF_IMG_20150614_125741_cell_101.png
- Uninfected/C130P91ThinF_IMG_20151004_142709_cell_101.png
- Uninfected/C131P92ThinF_IMG_20151004_145410_cell_3.png
- Uninfected/C136P97ThinF_IMG_20151005_141321_cell_15.png

- Uninfected/C137P98ThinF_IMG_20151005_160122_cell_29.png
- Uninfected/C138P99ThinF_IMG_20151005_172849_cell_90.png
- Uninfected/C13NThinF_IMG_20150614_131457_cell_193.png
- Uninfected/C141P102ThinF_IMG_20151005_214615_cell_132.png
- Uninfected/C142P103ThinF_IMG_20151005_221115_cell_128.png
- Uninfected/C150P111ThinF_IMG_20151115_120244_cell_108.png
- Uninfected/C150P111ThinF_IMG_20151115_120244_cell_187.png
- Uninfected/C151P112ThinF_IMG_20151115_121644_cell_131.png
- Uninfected/C152P113ThinF_IMG_20151115_124032_cell_171.png
- Uninfected/C153P114ThinF_IMG_20151115_135639_cell_39.png
- Uninfected/C157P118ThinF_IMG_20151115_163915_cell_107.png
- Uninfected/C158P119ThinF_IMG_20151115_181136_cell_91.png
- Uninfected/C169P130ThinF_IMG_20151118_163539_cell_51.png
- Uninfected/C169P130ThinF_IMG_20151118_173039_cell_245.png
- Uninfected/C188P149ThinF_IMG_20151203_135433_cell_88.png
- Uninfected/C205ThinF_IMG_20151106_151800_cell_72.png
- Uninfected/C206ThinF_IMG_20151029_140917_cell_192.png
- Uninfected/C207ThinF_IMG_20151029_143711_cell_252.png
- Uninfected/C207ThinF_IMG_20151029_143711_cell_6.png
- Uninfected/C208ThinF_IMG_20151029_155554_cell_149.png
- Uninfected/C210ThinF_IMG_20151029_162834_cell_57.png
- Uninfected/C210ThinF_IMG_20151029_162934_cell_258.png
- Uninfected/C214ThinF_IMG_20151106_115440_cell_163.png
- Uninfected/C214ThinF_IMG_20151106_115440_cell_58.png
- Uninfected/C214ThinF_IMG_20151106_131748_cell_148.png
- Uninfected/C216ThinF_IMG_20151106_135228_cell_220.png
- Uninfected/C216ThinF_IMG_20151106_135337_cell_60.png
- Uninfected/C216ThinF_IMG_20151106_135653_cell_218.png
- Uninfected/C225ThinF_IMG_20151112_113735_cell_81.png

- Uninfected/C226ThinF_IMG_20151112_131408_cell_102.png
- Uninfected/C228ThinF_IMG_20151112_142452_cell_181.png
- Uninfected/C229ThinF_IMG_20151112_144147_cell_154.png
- Uninfected/C231ThinF_IMG_20151112_153246_cell_134.png
- Uninfected/C239ThinF_IMG_20151127_113008_cell_154.png
- Uninfected/C37BP2_thinF_IMG_20150620_132847a_cell_37.png
- Uninfected/C37BP2_thinF_IMG_20150620_133238a_cell_41.png
- Uninfected/C38P3thinF_original_IMG_20150621_112116_cell_60.png
- Uninfected/C38P3thinF_original_IMG_20150621_112227_cell_11.png
- Uninfected/C39P4thinF_original_IMG_20150622_105102_cell_59.png
- Uninfected/C39P4thinF_original_IMG_20150622_111206_cell_22.png
- Uninfected/C48P9thinF_IMG_20150721_160944_cell_177.png
- Uninfected/C54P15thinF_IMG_20150728_112832_cell_24.png
- Uninfected/C55P16thinF_IMG_20150728_123237_cell_15.png
- Uninfected/C57P18thinF_IMG_20150729_110457_cell_78.png
- Uninfected/C5NThinF_IMG_20150609_122227_cell_9.png
- Uninfected/C60P21thinF_IMG_20150803_144003_cell_81.png
- Uninfected/C61P22N_ThinF_IMG_20150818_112252_cell_144.png
- Uninfected/C61P22N_ThinF_IMG_20150818_112811_cell_128.png
- Uninfected/C62P23N_ThinF_IMG_20150818_133211_cell_10.png
- Uninfected/C62P23N_ThinF_IMG_20150818_133211_cell_94.png
- Uninfected/C62P23N_ThinF_IMG_20150818_133307_cell_137.png
- Uninfected/C66P27N_ThinF_IMG_20150818_163419_cell_189.png
- Uninfected/C69P30N_ThinF_IMG_20150819_135421_cell_157.png
- Uninfected/C6NThinF_IMG_20150609_121955_cell_133.png
- Uninfected/C6NThinF_IMG_20150609_122725_cell_208.png
- Uninfected/C6NThinF_IMG_20150609_122832_cell_89.png
- Uninfected/C74P35_ThinF_IMG_20150815_120957_cell_11.png
- Uninfected/C75P36_ThinF_IMG_20150815_163707_cell_33.png

- Uninfected/C77P38ThinF_IMG_20150601_155125_cell_115.png
- Uninfected/C79P40ThinF_IMG_20150817_102702_cell_39.png
- Uninfected/C79P40ThinF_IMG_20150817_102702_cell_64.png
- Uninfected/C7NthinF_IMG_20150611_104609_cell_62.png
- Uninfected/C86P47ThinF_IMG_20150820_124943_cell_207.png
- Uninfected/C87P48ThinF_IMG_20150820_141913_cell_165.png
- Uninfected/C88P49ThinF_IMG_20150820_150724_cell_129.png
- Uninfected/C91P52ThinF_IMG_20150821_124504_cell_60.png
- Uninfected/C91P52ThinF_IMG_20150821_124937_cell_171.png
- Uninfected/C92P53ThinF_IMG_20150821_151646_cell_155.png
- Uninfected/C93P54ThinF_IMG_20150821_162901_cell_60.png
- Uninfected/C93P54ThinF_IMG_20150821_164024_cell_158.png
- Uninfected/C94P55ThinF_IMG_20150821_165519_cell_165.png
- Uninfected/C99P60ThinF_IMG_20150918_141001_cell_69.png
- Uninfected/C99P60ThinF_IMG_20150918_141129_cell_84.png
- Uninfected/C99P60ThinF_IMG_20150918_142128_cell_45.png

pneumonia

155 flagged identifiers:

- NORMAL/IM-0133-0001.jpeg
- NORMAL/IM-0621-0001.jpeg
- NORMAL/NORMAL2-IM-0587-0001-0002.jpeg
- NORMAL/NORMAL2-IM-0611-0001.jpeg
- NORMAL/NORMAL2-IM-0718-0001.jpeg
- NORMAL/NORMAL2-IM-0723-0001.jpeg
- NORMAL/NORMAL2-IM-0730-0001.jpeg
- NORMAL/NORMAL2-IM-0741-0001.jpeg
- NORMAL/NORMAL2-IM-0757-0001.jpeg
- NORMAL/NORMAL2-IM-0808-0001.jpeg
- NORMAL/NORMAL2-IM-1064-0001.jpeg

- NORMAL/NORMAL2-IM-1067-0001.jpeg
- NORMAL/NORMAL2-IM-1214-0001.jpeg
- NORMAL/NORMAL2-IM-1250-0001.jpeg
- NORMAL/NORMAL2-IM-1356-0001.jpeg
- NORMAL/NORMAL2-IM-1360-0001.jpeg
- NORMAL/NORMAL2-IM-1362-0001.jpeg
- NORMAL/NORMAL2-IM-1379-0001.jpeg
- NORMAL/NORMAL2-IM-1396-0001.jpeg
- NORMAL/NORMAL2-IM-1400-0001.jpeg
- NORMAL/NORMAL2-IM-1406-0001.jpeg
- NORMAL/NORMAL2-IM-1437-0001.jpeg
- NORMAL/NORMAL2-IM-1438-0001.jpeg
- PNEUMONIA/person1006_bacteria_2937.jpeg
- PNEUMONIA/person1049_bacteria_2983.jpeg
- PNEUMONIA/person104_virus_191.jpeg
- PNEUMONIA/person109_virus_203.jpeg
- PNEUMONIA/person10_bacteria_43.jpeg
- PNEUMONIA/person1107_bacteria_3048.jpeg
- PNEUMONIA/person1114_bacteria_3055.jpeg
- PNEUMONIA/person111_virus_212.jpeg
- PNEUMONIA/person1140_virus_1885.jpeg
- PNEUMONIA/person1149_virus_1925.jpeg
- PNEUMONIA/person1152_virus_1930.jpeg
- PNEUMONIA/person1156_virus_1935.jpeg
- PNEUMONIA/person1158_virus_1938.jpeg
- PNEUMONIA/person1160_virus_1947.jpeg
- PNEUMONIA/person1162_bacteria_3107.jpeg
- PNEUMONIA/person1162_virus_1950.jpeg
- PNEUMONIA/person1163_virus_1951.jpeg

- PNEUMONIA/person1164_virus_1958.jpeg
- PNEUMONIA/person1171_bacteria_3118.jpeg
- PNEUMONIA/person1174_virus_1980.jpeg
- PNEUMONIA/person1175_bacteria_3122.jpeg
- PNEUMONIA/person1180_virus_2015.jpeg
- PNEUMONIA/person1181_bacteria_3129.jpeg
- PNEUMONIA/person1183_bacteria_3131.jpeg
- PNEUMONIA/person1183_virus_2018.jpeg
- PNEUMONIA/person1184_virus_2019.jpeg
- PNEUMONIA/person1187_virus_2023.jpeg
- PNEUMONIA/person1190_virus_2031.jpeg
- PNEUMONIA/person1191_virus_2032.jpeg
- PNEUMONIA/person1208_bacteria_3160.jpeg
- PNEUMONIA/person1227_virus_2078.jpeg
- PNEUMONIA/person1244_bacteria_3200.jpeg
- PNEUMONIA/person1281_virus_2204.jpeg
- PNEUMONIA/person1287_virus_2210.jpeg
- PNEUMONIA/person12_bacteria_48.jpeg
- PNEUMONIA/person1357_virus_2338.jpeg
- PNEUMONIA/person1358_virus_2339.jpeg
- PNEUMONIA/person1368_virus_2353.jpeg
- PNEUMONIA/person1405_bacteria_3566.jpeg
- PNEUMONIA/person1405_virus_2408.jpeg
- PNEUMONIA/person1409_bacteria_3583.jpeg
- PNEUMONIA/person1449_bacteria_3746.jpeg
- PNEUMONIA/person1453_bacteria_3770.jpeg
- PNEUMONIA/person1454_bacteria_3782.jpeg
- PNEUMONIA/person1455_virus_2496.jpeg
- PNEUMONIA/person1459_bacteria_3797.jpeg







- PNEUMONIA/person1461_virus_2510.jpeg
- PNEUMONIA/person1475_virus_2558.jpeg
- PNEUMONIA/person1483_virus_2574.jpeg
- PNEUMONIA/person1490_virus_2596.jpeg
- PNEUMONIA/person20_bacteria_69.jpeg
- PNEUMONIA/person21_bacteria_73.jpeg
- PNEUMONIA/person275_virus_565.jpeg
- PNEUMONIA/person289_virus_593.jpeg
- PNEUMONIA/person293_bacteria_1379.jpeg
- PNEUMONIA/person294_bacteria_1386.jpeg
- PNEUMONIA/person294_virus_606.jpeg
- PNEUMONIA/person295_bacteria_1390.jpeg
- PNEUMONIA/person295_virus_612.jpeg
- PNEUMONIA/person299_bacteria_1419.jpeg
- PNEUMONIA/person301_bacteria_1429.jpeg
- PNEUMONIA/person302_bacteria_1430.jpeg
- PNEUMONIA/person303_bacteria_1431.jpeg
- PNEUMONIA/person311_bacteria_1453.jpeg
- PNEUMONIA/person317_bacteria_1473.jpeg
- PNEUMONIA/person327_bacteria_1509.jpeg
- PNEUMONIA/person343_bacteria_1584.jpeg
- PNEUMONIA/person344_bacteria_1585.jpeg
- PNEUMONIA/person357_bacteria_1639.jpeg
- PNEUMONIA/person359_bacteria_1642.jpeg
- PNEUMONIA/person384_bacteria_1755.jpeg
- PNEUMONIA/person38_bacteria_196.jpeg
- PNEUMONIA/person391_bacteria_1782.jpeg
- PNEUMONIA/person409_bacteria_1824.jpeg
- PNEUMONIA/person410_bacteria_1825.jpeg

- PNEUMONIA/person411_bacteria_1826.jpeg
- PNEUMONIA/person416_bacteria_1840.jpeg
- PNEUMONIA/person422_bacteria_1853.jpeg
- PNEUMONIA/person423_virus_869.jpeg
- PNEUMONIA/person432_virus_881.jpeg
- PNEUMONIA/person444_bacteria_1927.jpeg
- PNEUMONIA/person450_bacteria_1941.jpeg
- PNEUMONIA/person454_virus_938.jpeg
- PNEUMONIA/person458_virus_945.jpeg
- PNEUMONIA/person45_bacteria_222.jpeg
- PNEUMONIA/person46_bacteria_225.jpeg
- PNEUMONIA/person472_virus_969.jpeg
- PNEUMONIA/person476_bacteria_2026.jpeg
- PNEUMONIA/person477_bacteria_2031.jpeg
- PNEUMONIA/person488_bacteria_2062.jpeg
- PNEUMONIA/person489_bacteria_2067.jpeg
- PNEUMONIA/person491_bacteria_2071.jpeg
- PNEUMONIA/person492_bacteria_2085.jpeg
- PNEUMONIA/person493_bacteria_2087.jpeg
- PNEUMONIA/person494_bacteria_2088.jpeg
- PNEUMONIA/person496_bacteria_2095.jpeg
- PNEUMONIA/person504_bacteria_2127.jpeg
- PNEUMONIA/person504_bacteria_2133.jpeg
- PNEUMONIA/person506_bacteria_2138.jpeg
- PNEUMONIA/person534_bacteria_2251.jpeg
- PNEUMONIA/person536_virus_1064.jpeg
- PNEUMONIA/person545_bacteria_2290.jpeg
- PNEUMONIA/person549_bacteria_2303.jpeg
- PNEUMONIA/person575_bacteria_2374.jpeg

- PNEUMONIA/person58_bacteria_278.jpeg
- PNEUMONIA/person590_bacteria_2428.jpeg
- PNEUMONIA/person610_bacteria_2475.jpeg
- PNEUMONIA/person618_bacteria_2489.jpeg
- PNEUMONIA/person630_bacteria_2512.jpeg
- PNEUMONIA/person645_bacteria_2537.jpeg
- PNEUMONIA/person652_bacteria_2544.jpeg
- PNEUMONIA/person674_bacteria_2568.jpeg
- PNEUMONIA/person68_bacteria_336.jpeg
- PNEUMONIA/person703_virus_1300.jpeg
- PNEUMONIA/person72_bacteria_354.jpeg
- PNEUMONIA/person732_virus_1353.jpeg
- PNEUMONIA/person734_bacteria_2637.jpeg
- PNEUMONIA/person745_bacteria_2648.jpeg
- PNEUMONIA/person74_bacteria_363.jpeg
- PNEUMONIA/person781_virus_1412.jpeg
- PNEUMONIA/person788_virus_1419.jpeg
- PNEUMONIA/person802_bacteria_2708.jpeg
- PNEUMONIA/person805_bacteria_2712.jpeg
- PNEUMONIA/person813_virus_1449.jpeg
- PNEUMONIA/person815_bacteria_2726.jpeg
- PNEUMONIA/person833_virus_1469.jpeg
- PNEUMONIA/person844_virus_1487.jpeg
- PNEUMONIA/person86_virus_159.jpeg
- PNEUMONIA/person885_bacteria_2809.jpeg
- PNEUMONIA/person8_bacteria_37.jpeg
- PNEUMONIA/person929_bacteria_2854.jpeg
- PNEUMONIA/person94_virus_176.jpeg

References

- [1] Rajaraman S. et al. Pre-trained convolutional neural networks as feature extractors toward improved malaria parasite detection in thin blood smear images. *PeerJ* 6:e4568, 2018.
- [2] Kermany D. et al. Identifying medical diagnoses and treatable diseases by image-based deep learning. *Cell* 172(5):1122–1131, 2018.
- [3] Kipf T., Welling M. Semi-supervised classification with graph convolutional networks. *ICLR*, 2017.
- [4] Northcutt C., Jiang L., Chuang I. Confident learning: estimating uncertainty in dataset labels. *JAIR* 70:1373–1411, 2021.
- [5] Kingma D., Ba J. Adam: a method for stochastic optimization. *ICLR*, 2015.






Figure 7: Pneumonia confident joint from the 3-fold out-of-fold census. The off-diagonal mass is again strongly one-directional: 973 training images labelled *pneumonia* cross-validate as *normal*, vs 23 in the opposite direction.




Figure 8: Out-of-fold confidence by true label (seaborn KDE). The one-directional label-noise mass is visible as the heavy wrong-side tails: Parasitized-labelled cells with near-zero parasite confidence, and PNEUMONIA-labelled films with near-zero pneumonia confidence.